

AI × Security Landscape

Initial survey, themes, exits, and ideas to pursue

Prepared 2026-04-20 | For Zubair Ashraf

Contents

1. Executive summary

This report captures a comprehensive initial survey of startups at the intersection of AI and cybersecurity. It covers 83 active companies across 12 sub-categories, summarizes notable funding rounds and exits from 2024 through Q1 2026, distills recurring themes and contrarian approaches, and proposes fifteen new-venture ideas. A companion tracker (AI_Security_Startup_Tracker.xlsx) contains the structured data and a weekly log that will be appended every Monday.

Three events have defined the 2024-2026 AI-security cycle. First, the unprecedented \$32B agreement for Google to acquire Wiz in March 2025 reset the bar for platform consolidation. Second, Palo Alto Networks' acquisition of Protect AI (closed July 2025) and Cato Networks' first-ever acquisition of Aim Security (September 2025) signaled that platform vendors will buy rather than build AI-specific security. Third, Cyera's \$540M Series E at a \$6B valuation (June 2025) confirmed that data security for AI is the breakout category.

On the offensive side, purpose-built AI red teams and autonomous pen-testers (Haize Labs, Grey Swan, XBOW, SplxAI) are scaling alongside AI-SOC analysts (Dropzone, Prophet, Radiant, Simbian, Crogl). The agentic security layer is emerging but has not yet produced a durable category leader.

2. Methodology

Companies were selected if they (a) apply AI or ML as a first-class capability to a security problem, or (b) provide security for AI systems themselves. Pure-play AI infrastructure vendors without a security angle and pure-play security vendors without an AI angle were excluded.

Funding figures are drawn from company press releases, investor announcements, and tier-one reporting (Reuters, TechCrunch, SC World, CRN). Where the public record is thin, the tracker marks values as "n/d" (not disclosed). Valuations are post-money unless noted. The data is refreshed weekly via a scheduled task.

3. Market landscape

We organized the market into twelve sub-categories. The table below lists each theme, its thesis, and representative players. Detailed one-line descriptions per company are in section 4.

Theme	Description	Representative players
AI-SOC / autonomous SOC analyst	LLM-based agents triage, enrich, and investigate alerts end-to-end, reducing Tier-1/2 headcount. The category exploded in 2024 and has matured in 2025-2026 as early adopters (Microsoft Security Copilot, CrowdStrike Charlotte, Splunk AI) shipped native options. The independent players differentiate on SIEM-agnostic reasoning, custom integrations, and MSSP-friendly deployments.	Dropzone AI (\$37M Series B, Jul 2025), Prophet Security, Radiant Security, Simbian, Crogl, Exaforce, Seven.AI, Andesite AI, Torq (unicorn 2025), Tines, Hunters, Anvilogic.
GenAI firewall / prompt security	Inline guardrails sit between end users and LLM applications to block prompt injection, jailbreaks, PII disclosure, and policy violations. Standardization is emerging around the OWASP LLM Top 10 and NIST AI RMF. Consolidation has already begun: Cisco acquired Robust Intelligence (Aug 2024), and Palo Alto Networks	Lakera, Prompt Security, Apex Security, Pillar Security, CalypsoAI, WitnessAI, Noma Security, HiddenLayer, Protect AI (PANW), Robust

Theme	Description	Representative players
	closed Protect AI (Jul 2025).	Intelligence (Cisco).
AI red teaming / AI pentest	Automated adversarial evaluation of models, agents, and AI apps, ranging from continuous jailbreak generation to full autonomous pen-testing of production systems. Services-heavy today; the bull case is productized continuous testing as EU AI Act / NIST AI RMF obligations mature.	Haize Labs, Grey Swan AI, Mindgard, SplxAI, Repello, XBOW (\$75M Sequoia-led Series A), ZeroPath.
Agentic AI security	Runtime policy, tool-use guardrails, and identity for autonomous AI agents. This is the most speculative category — definitions are still fluid — but the biggest long-term prize if agents become the dominant application paradigm. Expect category overlap with NHI and GenAI firewall.	Noma Security, Straiker, Invariant Labs, Conifers.ai, Simbian, Andesite.
Non-human identity / machine identity	Discovery, lifecycle, and ITDR for service accounts, API keys, service principals, and now AI agents. Breach data (Internet Archive, Snowflake, US Treasury OCC) validated the thesis. CyberArk's Venafi acquisition and ServiceNow's Veza deal signal heavy incumbent interest.	Astrix (\$45M Series B, Sep 2025), Oasis Security, Token Security, Clutch Security, Entro Security, Venafi (now CyberArk), Veza (ServiceNow, pending).
Data security for AI (DSPM, AI-DLP)	Discover, classify, and govern sensitive data that flows to and from AI systems. The fastest-scaling category of the past 24 months, anchored by Cyera's \$6B valuation and \$540M Series E (Jun 2025). Expected to partially consolidate into CNAPPs.	Cyera, Sentra, Normalyze, BigID, Cyberhaven, Harmonic Security, Securi AI, Opsin.
AI governance / model risk	Governance, AI-BOM, policy, and audit aligned with EU AI Act and NIST AI RMF. Growing as procurement line-item in 2026 budgets. Software margins are lower than security norm; opportunity is in enterprise data fabric + continuous monitoring.	Credo AI, Holistic AI, Cranium, Monitaur, Trustible, Fiddler AI, Arthur AI.
Deepfake / synthetic media defense	Detection and provenance for AI-generated voice, video, and image content. Deepfake-enabled CFO vishing against enterprises accelerated through 2025, and attention has shifted from detection-only to provenance (C2PA) and KYC hardening.	Adaptive Security (\$43M Series A; a16z + OpenAI Startup Fund), Reality Defender, GetReal Labs, Clarity.
AI for code security	LLMs applied to vulnerability discovery, auto-fix, secure coding, and SCA. GenAI enables embedded AppSec directly in IDEs and CI, a long-time industry goal. Copilot-native defenses (GitHub Advanced Security, Snyk DeepCode) are the primary competitive pressure.	Semgrep, Socket, Endor Labs, Apiiro, Aikido, Corga, Pixee, Ghost Security, XBOW.
AI-native email & social engineering	Behavioral detection of BEC, phishing, QR-phishing, voice deepfake attacks, and collaboration tool abuse. Abnormal's Series D at \$5.1B (Aug 2024) and IPO track established that a standalone AI-first email-security vendor can scale to near-public.	Abnormal Security, Material Security, Adaptive Security, Tessian (now Proofpoint), Armorblox (Cisco).
AI-enabled GRC / compliance	Policy-as-code, evidence collection, and continuous control monitoring with LLM summarization. Vanta, Drata, and Secureframe have added AI assistants; differentiation is in depth of control coverage (NIS2, DORA, EU AI Act, HIPAA) and quality of evidence generation.	Vanta (Series C at \$2.45B), Drata, Secureframe, Anecdotes, Zenity (copilot GRC).
Cloud + runtime AI	AI-driven detection of cloud workloads and container	Wiz (pending Google deal),

Theme	Description	Representative players
	runtime anomalies. The Wiz / Google announcement (\$32B) is the tectonic event of 2025 — a reshaping of CNAPP and the broader security platform stack that is still propagating through the market.	Orca Security, Upwind, Sweet Security, RAD Security, Kodem.

4. Companies by category

83 companies across the twelve themes. The table includes headquarters, current stage, cumulative disclosed funding (or latest valuation), and a one-line product angle. For full data, see the Companies tab in the tracker.

AI-SOC / autonomous SOC analyst

Company	HQ	Stage	Funding / Valuation	Angle
Dropzone AI	Seattle	Series B	\$60M total	Autonomous AI SOC analyst.
Prophet Security	San Francisco	Series A	\$30M	Ex-Netskope team; noisy-alert reasoning.
Radiant Security	San Francisco	Series A	\$15M	Adaptive investigation copilot.
Simbian	Mountain View	Seed+	\$15M	Agentic SOC with guardrails.
Crogl	San Francisco	Seed+	\$25M	Investigation copilot; ex-Databricks.
Exaforce	San Francisco	Seed	\$7M	Ex-Prisma Cloud leadership.
Seven.AI	Tel Aviv	Seed	\$6M	Ex-Check Point; AI investigator.
Andesite AI	Cambridge MA	Series A	\$19M	Threat detection with AI reasoning.
Conifers.ai	Tel Aviv	Seed	\$10M	Multi-agent SOC orchestration.
Torq	Tel Aviv	Series C	\$192M (\$1B val)	AI hyperautomation for SOC.
Tines	Dublin	Series C	\$146M (\$1.13B val)	Workflow automation with AI agents.
Hunters	Tel Aviv	Series C	\$118M	AI-driven SIEM.
Anvilogic	Palo Alto	Series B	\$45M	AI detection engineering.
SnapAttack	McLean	Series A	\$25M	Generative detection authoring.
Query.AI	Boulder	Series A	\$19M	Federated AI search.
Nightfall AI	San Francisco	Series B	\$59M	AI DLP/CASB.
Salem Cyber	New York	Seed	\$4M	AI alert triage for MSSPs.

GenAI firewall / prompt security

Company	HQ	Stage	Funding / Valuation	Angle
Lakera	Zurich	Series A	\$30M	Real-time LLM guardrails.
HiddenLayer	Austin	Series A	\$50M	Model integrity + AI-BOM.
Protect AI	Seattle	Acquired (PANW)	\$108M	End-to-end MLSecOps; exited Jul 2025.
Robust Intelligence	Berkeley	Acquired (Cisco)	\$44M	AI firewall; exited Aug 2024.
Prompt Security	Tel Aviv	Series A	\$23M	Inline enterprise GenAI guardrail.
CalypsoAI	Dublin / Virginia	Series A	\$38M	Model-agnostic scanners.
WitnessAI	Berkeley	Series A	\$27M	Intent-based policy for employee GenAI.
Pillar Security	Tel Aviv	Seed	\$9M	AI lifecycle security.
Apex Security	Tel Aviv	Seed	\$7M	GenAI app security.
Cranium	Short Hills NJ	Series A	\$32M	AI supply chain + AI-BOM.
Aim Security	Tel Aviv	Acquired (Cato)	\$28M	Workplace AI governance; exited Sep 2025.

AI red teaming / AI pentest

Company	HQ	Stage	Funding / Valuation	Angle
Haize Labs	New York	Seed	\$5M	Automated jailbreak generation.
Grey Swan AI	New York	Seed+	\$15M	Circuit-breaker defenses + red team.
Mindgard	London	Seed+	\$10M	Adversarial testing for ML+LLMs.
SplxAI	Zagreb	Seed+	\$7M	Automated pentest + remediation for GenAI.
XBOW	London	Series A	\$75M	Autonomous AI pentester (Semmler alumni).
ZeroPath	San Francisco	Seed	\$5M	AI-native SAST with exploit reasoning.
Repello AI	Bangalore	Seed	\$2M	Continuous GenAI red teaming.
Promptfoo	San Francisco	Seed	\$5M	LLM eval + red team tooling.

Agentic AI security

Company	HQ	Stage	Funding / Valuation	Angle
Noma Security	Tel Aviv / NYC	Series B	~\$132M	AI/data pipeline security; agent runtime.
Straiker	Palo Alto	Seed	\$8M	Runtime AI agent defense.
Invariant Labs	Zurich	Seed	\$7M	Policy engine for agent tool use.

Non-human / machine identity

Company	HQ	Stage	Funding / Valuation	Angle
Astrix Security	Tel Aviv	Series B	\$85M (\$400M val)	NHI security + agent identity.
Oasis Security	New York	Series A	\$40M	NHI discovery & lifecycle.
Token Security	Tel Aviv	Series A	\$27M	Machine ITDR.
Clutch Security	New York	Series A	\$28M	Universal NHI platform.
Entro Security	Boston	Series A	\$24M	Secrets discovery and monitoring.

Data security for AI (DSPM)

Company	HQ	Stage	Funding / Valuation	Angle
Cyera	New York / Tel Aviv	Series E	\$1.3B (\$6B val)	AI-powered DSPM.
Sentra	New York / Tel Aviv	Series C	\$110M	DSPM with native AI classifier.
Normalyze	Los Altos	Series A	\$27M	Agentless DSPM.
Theom	Palo Alto	Seed	\$4M	Data risk quantification.
BigID	New York	Growth	\$315M (\$1.25B val)	ML-driven data discovery.
Cyberhaven	San Francisco	Series D	\$213M (\$1B val)	Data-lineage DLP.
Harmonic Security	London	Series A	\$17.5M	GenAI data leakage prevention.
Securiti AI	San Jose	Series C	\$156M	DataCommand Center.
Opsin Security	Austin	Seed	\$6M	AI data access governance.

AI governance / model risk

Company	HQ	Stage	Funding / Valuation	Angle
Credo AI	San Francisco	Series A	\$31M	Enterprise AI governance.
Fiddler AI	Palo Alto	Series C	\$77M	ML + LLM monitoring.
Arthur AI	New York	Series B	\$60M	Model monitoring + Arthur Shield LLM firewall.
Arize AI	Berkeley	Series B	\$61M	LLM eval, tracing, guardrails.
Galileo	San Francisco	Series B	\$68M	LLM eval + Protect guardrails.
Holistic AI	London	Series A	\$23M	Audit for regulated industries.
Zenity	Tel Aviv	Series B	\$55M	Copilot / low-code security.

Deepfake / synthetic media

Company	HQ	Stage	Funding / Valuation	Angle
Adaptive Security	New York	Series A	\$58M (\$500M val)	Deepfake-aware phishing defense (a16z + OpenAI).
Reality Defender	New York	Series A	\$33M	Multi-modal deepfake detection.
GetReal Labs	San Francisco	Seed	\$17.5M	Hany Farid-led; forensic and C2PA.
Clarity	Tel Aviv	Series A	\$16M	Deepfake + voice AI defense.

AI for code / app security

Company	HQ	Stage	Funding / Valuation	Angle
Semgrep	San Francisco	Series D	\$151M (\$500M val)	AI remediation + SAST.
Socket	San Francisco	Series B	\$60M	Supply-chain AI.
Endor Labs	Palo Alto	Series B	\$163M (\$800M val)	Reachability-based SCA + AI.
Apiiro	Tel Aviv / NYC	Series C	\$135M (\$1.1B val)	ASPM with AI.
Chainguard	Kirkland	Series D	\$256M (\$3.5B val)	Hardened containers.
Aikido Security	Ghent	Series A	\$17M	All-in-one AppSec for SMBs.
Ghost Security	Austin	Series A	\$30M	Runtime AI for APIs.
Corgea	San Francisco	Seed	\$5M	Auto-fix SAST findings with LLMs.

AI-native email / BEC

Company	HQ	Stage	Funding / Valuation	Angle
Abnormal Security	San Francisco	Series D	\$546M (\$5.1B val)	Behavioral AI email; IPO track.
Material Security	Redwood City	Series C	\$115M (\$1.1B val)	Email + workspace security.

AI-enabled GRC / compliance

Company	HQ	Stage	Funding / Valuation	Angle
Vanta	San Francisco	Series C	\$354M (\$2.45B val)	Compliance automation; Vanta AI.
Drata	San Diego	Series C	\$328M (\$2B val)	Continuous control monitoring.
Secureframe	San Francisco	Series C	\$79M	AI-generated evidence.
Anecdotes	Palo Alto	Series B	\$55M	GRC data fabric.

Cloud + runtime AI

Company	HQ	Stage	Funding / Valuation	Angle
Wiz	NYC / Tel Aviv	Pending acq. (Google)	\$1.9B raised; \$32B deal	Cloud security platform.
Orca Security	Portland	Series C	\$632M (\$1.8B val)	Agentless cloud security.
Upwind	Tel Aviv	Series B	\$180M	Runtime-first CNAPP with AI.
Sweet Security	Tel Aviv	Series A	\$45M	Cloud runtime AI detection.
RAD Security	San Francisco	Series A	\$22M	Behavioral cloud workload fingerprinting.

5. Funding activity, 2024 - Q1 2026

Selected rounds ≥ \$30M or of strategic significance. Complete list in tracker.

Date	Company	Round	Amount	Valuation	Category
2026-03	Cyera	Secondary/Ext	\$150M	\$7B	DSPM
2025-11	Abnormal Security	Pre-IPO / Secondary	\$200M	\$6.5B	AI email
2025-09	Noma Security	Series B	\$100M	n/d	AI lifecycle
2025-09	Astrix Security	Series B	\$45M	\$400M	NHI
2025-07	Dropzone AI	Series B	\$37M	n/d	AI-SOC
2025-06	Cyera	Series E	\$540M	\$6B	DSPM
2025-06	XBOW	Series A	\$75M	n/d	AI pentest
2025-05	Chainguard	Series D	\$140M	\$3.5B	Containers
2025-04	Wiz (announced)	Acquisition	\$32B	\$32B	CNAPP
2025-04	Adaptive Security	Series A	\$43M	\$500M	Deepfake
2025-03	Semgrep	Series D	\$100M	\$500M	AppSec AI
2025-02	Upwind	Series B	\$100M	\$900M	CNAPP
2025-02	Push Security	Series B	\$30M	n/d	Identity
2025-01	Apiiro	Series C	\$60M	\$1.1B	ASPM
2024-11	Cyera	Series D	\$300M	\$3B	DSPM
2024-10	Protect AI	Series B	\$60M	\$400M	MLSecOps
2024-09	Endor Labs	Series B	\$70M	\$800M	SCA AI
2024-08	Abnormal Security	Series D	\$250M	\$5.1B	AI email
2024-07	HiddenLayer	Series A	\$50M	n/d	ML security
2024-06	Pentera	Series D	\$60M	\$1B	Pentest
2024-05	Torq	Series C	\$70M	\$1B	SecOps automation
2024-05	Wiz	Series E	\$1B	\$12B	CNAPP
2024-03	Vanta	Series C	\$150M	\$2.45B	GRC

6. Exits & M&A

The 2024-2026 window produced the largest security M&A events on record, including the pending \$32B Google/Wiz deal and \$28B Cisco/Splunk close. Below is the consolidated exit ledger.

Date	Deal	Value	Category
2025-12	Veza → ServiceNow (pending)	n/d	Identity security
2025-09	Aim Security → Cato Networks	n/d	AI workplace security
2025-07	Protect AI → Palo Alto Networks	~\$500-700M est.	MLSecOps
2025-03	Wiz → Google (pending)	\$32B	CNAPP
2024-11	Own → Salesforce	\$1.9B	SaaS data protection
2024-10	Darktrace → Thoma Bravo	\$5.3B	AI threat detection
2024-09	Recorded Future → Mastercard	\$2.65B	Threat intel
2024-08	Robust Intelligence → Cisco	n/d	AI firewall
2024-06	Venafi → CyberArk	\$1.54B	Machine identity
2024-03	Splunk → Cisco	\$28B	SIEM
2024-02	Lacework → Fortinet	n/d (distressed)	CNAPP
2023-11	Dig Security → Palo Alto	\$400M	DSPM
2023-07	Laminar → Rubrik	\$250M	DSPM
2023-07	Armorblox → Cisco	n/d	AI email

7. Themes in depth

AI-SOC / autonomous SOC analyst

LLM-based agents triage, enrich, and investigate alerts end-to-end, reducing Tier-1/2 headcount. The category exploded in 2024 and has matured in 2025-2026 as early adopters (Microsoft Security Copilot, CrowdStrike Charlotte, Splunk AI) shipped native options. The independent players differentiate on SIEM-agnostic reasoning, custom integrations, and MSSP-friendly deployments.

Representative players: Dropzone AI (\$37M Series B, Jul 2025), Prophet Security, Radiant Security, Simbian, Crogl, Exaforce, Seven.AI, Andesite AI, Torq (unicorn 2025), Tines, Hunters, Anvillogic.

GenAI firewall / prompt security

Inline guardrails sit between end users and LLM applications to block prompt injection, jailbreaks, PII disclosure, and policy violations. Standardization is emerging around the OWASP LLM Top 10 and NIST AI RMF. Consolidation has already begun: Cisco acquired Robust Intelligence (Aug 2024), and Palo Alto Networks closed Protect AI (Jul 2025).

Representative players: Lakera, Prompt Security, Apex Security, Pillar Security, CalypsoAI, WitnessAI, Noma Security, HiddenLayer, Protect AI (PANW), Robust Intelligence (Cisco).

AI red teaming / AI pentest

Automated adversarial evaluation of models, agents, and AI apps, ranging from continuous jailbreak generation to full autonomous pen-testing of production systems. Services-heavy today; the bull case is productized continuous testing as EU AI Act / NIST AI RMF obligations mature.

Representative players: Haize Labs, Grey Swan AI, Mindgard, SplxAI, Repello, XBOW (\$75M Sequoia-led Series A), ZeroPath.

Agentic AI security

Runtime policy, tool-use guardrails, and identity for autonomous AI agents. This is the most speculative category — definitions are still fluid — but the biggest long-term prize if agents become the dominant application paradigm. Expect category overlap with NHI and GenAI firewall.

Representative players: Noma Security, Straiker, Invariant Labs, Conifers.ai, Simbian, Andesite.

Non-human identity / machine identity

Discovery, lifecycle, and ITDR for service accounts, API keys, service principals, and now AI agents. Breach data (Internet Archive, Snowflake, US Treasury OCC) validated the thesis. CyberArk's Venafi acquisition and ServiceNow's Veza deal signal heavy incumbent interest.

Representative players: Astrix (\$45M Series B, Sep 2025), Oasis Security, Token Security, Clutch Security, Entro Security, Venafi (now CyberArk), Veza (ServiceNow, pending).

Data security for AI (DSPM, AI-DLP)

Discover, classify, and govern sensitive data that flows to and from AI systems. The fastest-scaling category of the past 24 months, anchored by Cyera's \$6B valuation and \$540M Series E (Jun 2025). Expected to partially consolidate into CNAPPs.

Representative players: Cyera, Sentra, Normalyze, BigID, Cyberhaven, Harmonic Security, Securiti AI, Opsin.

AI governance / model risk

Governance, AI-BOM, policy, and audit aligned with EU AI Act and NIST AI RMF. Growing as procurement line-item in 2026 budgets. Software margins are lower than security norm; opportunity is in enterprise data fabric + continuous monitoring.

Representative players: Credo AI, Holistic AI, Cranium, Monitaur, Trustible, Fiddler AI, Arthur AI.

Deepfake / synthetic media defense

Detection and provenance for AI-generated voice, video, and image content. Deepfake-enabled CFO vishing against enterprises accelerated through 2025, and attention has shifted from detection-only to provenance (C2PA) and KYC hardening.

Representative players: Adaptive Security (\$43M Series A; a16z + OpenAI Startup Fund), Reality Defender, GetReal Labs, Clarity.

AI for code security

LLMs applied to vulnerability discovery, auto-fix, secure coding, and SCA. GenAI enables embedded AppSec directly in IDEs and CI, a long-time industry goal. Copilot-native defenses (GitHub Advanced Security, Snyk DeepCode) are the primary competitive pressure.

Representative players: Semgrep, Socket, Endor Labs, Apiiro, Aikido, Corgea, Pixee, Ghost Security, XBOW.

AI-native email & social engineering

Behavioral detection of BEC, phishing, QR-phishing, voice deepfake attacks, and collaboration tool abuse. Abnormal's Series D at \$5.1B (Aug 2024) and IPO track established that a standalone AI-first email-security vendor can scale to near-public.

Representative players: Abnormal Security, Material Security, Adaptive Security, Tessian (now Proofpoint), Armorblox (Cisco).

AI-enabled GRC / compliance

Policy-as-code, evidence collection, and continuous control monitoring with LLM summarization. Vanta, Drata, and Secureframe have added AI assistants; differentiation is in depth of control coverage (NIS2, DORA, EU AI Act, HIPAA) and quality of evidence generation.

Representative players: Vanta (Series C at \$2.45B), Drata, Secureframe, Anecdotes, Zenity (copilot GRC).

Cloud + runtime AI

AI-driven detection of cloud workloads and container runtime anomalies. The Wiz / Google announcement (\$32B) is the tectonic event of 2025 — a reshaping of CNAPP and the broader security platform stack that is still propagating through the market.

Representative players: Wiz (pending Google deal), Orca Security, Upwind, Sweet Security, RAD Security, Kodem.

8. Unique approaches and differentiators

- Reasoning-first vs. rule-first SOC: Dropzone and Prophet take an agent-reasoning approach with minimal pre-built playbooks, while Simbian and Andesite combine rule graphs with LLM reasoning. Torq and Tines extend legacy SOAR with AI-native actions.
- Inline vs. air-gapped guardrails: Prompt Security and Lakera emphasize inline guardrails for live traffic. Protect AI (now PANW) and Cranium lean heavier on AI-BOM, scanning, and governance workflows.
- Data-first vs. identity-first security for AI: Cyera and Sentra argue the control plane is data. Astrix, Oasis, Token, and Clutch argue it is non-human identity. Both will coexist, but consolidation into platform vendors is expected.
- Autonomous pentest vs. continuous red team: Horizon3 and XBOW aim for full-autonomy adversaries against production targets. Haize Labs, Grey Swan, and SplxAI focus on continuous adversarial evaluation of models.
- Detection vs. provenance for deepfakes: Reality Defender and Clarity chase better detection. GetReal (Hany Farid) bets on C2PA provenance becoming the standard. Adaptive Security wraps detection into social-engineering training and response.
- Open-core vs. closed SaaS: Semgrep, Arize (Phoenix), Promptfoo, Socket, Endor, and Chainguard have sizable OSS footprints that drive enterprise deals. Closed SaaS incumbents like Abnormal and Material are the counter-strategy.
- Platform absorption vs. specialist: Platform vendors (CrowdStrike, Palo Alto, Cisco, Microsoft) are bundling AI-security into existing suites. Specialists must prove that an independent category wedge survives.

9. Ideas to pursue

1. Agentic IAM / delegated authorization for AI agents

Today's IAM (Okta, Entra) was designed for humans and service accounts. AI agents need scoped, delegated, auditable capabilities per tool call, with step-up approvals and consent ceremonies. Wedge is to deeply integrate with agent frameworks (LangGraph, CrewAI, AutoGen) and offer a policy engine that reasons about capabilities, not just APIs.

2. AI-native insider risk

When employees hand work to AI copilots and agents, the copilot's actions become the attack surface. Classic IRM vendors (DTEX, Forcepoint) don't model agent context or lineage. Build an IRM that unifies user activity, copilot prompts, agent tool calls, and data lineage to score risk in real time.

3. Deepfake defense for call centers and banking ops

Voice vishing of CFOs and wire-approval desks is a weekly headline. KYC and liveness vendors fixed onboarding but not ongoing voice channels. Opportunity: a liveness + behavioral fusion layer that plugs into Genesys, NICE, Five9, and Zoom Contact Center, with an enterprise control plane.

4. Security copilot for SWE agents

Devin-class SWE agents execute shells, open PRs, and touch production systems. Enterprises will not deploy them without a guardrail layer: policy for allowed commands, code classification, PR review, and secret-scoping. Distribute as VS Code extension + GitHub App + CI plugin.

5. AI supply-chain posture (AI-BOM + provenance)

Almost no enterprise can answer 'what model, dataset, and fine-tune is in production, and how was it built.' EU AI Act Article 50 and US sector regs will force this. Build AI-BOM + provenance (C2PA) with a continuous-monitoring layer.

6. AI-native DFIR

IR playbooks remain manual; breaches like MGM, Change Healthcare, and Snowflake revealed how costly this is. Build a DFIR platform that reconstructs attacker TTPs via AI reasoning over telemetry, offers on-call AI responders, and runs as an IR-retainer SaaS with a services overlay.

7. Governed retrieval for enterprise AI

RAG pipelines routinely leak entitlements because document ACLs are stripped during vectorization. Build an ACL-aware retrieval layer and policy engine that plugs into Azure AI Search, OpenSearch, Pinecone, Weaviate, and the major Copilot surfaces.

8. Continuous AI red team as a service

Snapshot assessments are obsolete for systems whose prompts and models change weekly. Deliver continuous red teaming combining a prompt corpus, model of the app, eval harness, and attack agents. Sell into regulated industries with EU AI Act reporting obligations.

9. Small-business AI phishing defense

Abnormal and Material are enterprise-priced. With AI-phish attacking SMBs, MSPs need a defense they can bundle. Build MSP-first phishing defense with per-seat pricing, native Microsoft 365 / Google Workspace coverage, and a QR/SMS/voice layer.

10. Agent-aware secrets vending

Agents read repos and CI. Their blast radius is bigger than a human developer's. Build a secrets vending layer that issues short-lived, scoped credentials per agent invocation, with detection when an agent deviates from its declared scope.

11. EU AI Act compliance automation

GPAI model risk classification, documentation, post-market monitoring, and serious-incident reporting become mandatory in 2026. Build a platform with prebuilt taxonomies, cross-framework mappings (NIST AI RMF, ISO 42001, EU AI Act), and continuous evidence generation.

12. Browser-based AI employee DLP

Employees paste secrets into ChatGPT, Claude, and Gemini. Network DLP misses in-browser data, and CASBs trail the long tail of AI apps. Build a managed browser + Chrome/Edge extension + LLM classifier that covers the inference channel.

13. Synthetic-media KYC and liveness

Incode, Jumio, and Onfido were built before voice/video injection attacks matured. Build sensor-fusion liveness (camera, accelerometer, microphone, active challenges) and a provenance backbone for onboarding, SIM swap, and 3DS flows.

14. AI threat-intel fusion

CTI products remain document-centric. Build an AI layer that fuses OSINT, dark-web, vuln data, and industry feeds into per-customer narratives, with agents that refine queries and produce briefs. Partner with IR firms and MSSPs for distribution.

15. Policy-as-code for AI agents

Rego/OPA was not designed for agent decision paths. Create a policy runtime that reasons about agent plans, tool calls, and data lineage, with simulation, a 'deny-explain' UX, and auditing built in.

10. Update cadence

Weekly Update — Week of 2026-07-13

Normal 7-day weekly run (window 2026-07-10 to 2026-07-17, attempt 1). Not a quarterly deep dive (the last, 2026-05-18, is now 60 days old — next due ~2026-08-16) and not an annual sweep (the last was 2026-06-26). The week's signal was concentrated in identity and infrastructure: Oak's \$60M stealth exit for an AI-native identity operating system and Valarian's \$50M Series A for sovereign AI infrastructure were the standout rounds, alongside First Recon AI's runtime launch and QIZ Security's post-quantum seed. Ten stub rows were remediated and one material correction was made — Cyera's June Series G is reconciled to \$600M at a \$12B valuation (the company row had been stale at \$6B).

Funding Activity

2026-07-15 — Oak: \$60M seed co-led by Accel, CRV and Greylock (with AlphaDrive Ventures, Hetz Ventures and angels), emerging from stealth to build an “AI-native Identity Operating System” — a single control plane governing every identity across an organization, whether human, machine or AI agent. Founded by Shai Morag, the serial founder behind Ermetic (sold to Tenable for \$265M) and Secdo (sold to Palo Alto Networks), with CPO Tal Marom (ex-Tenable, Salesforce). The product is already GA with unnamed enterprise customers and a ~50-person team built in stealth. A \$60M seed for a first-time-public company is an unusually large signal of conviction in the agent-identity control-plane thesis; the wedge is an AI connector framework that maps granted access to actual app usage and revokes stale permissions in real time rather than during periodic reviews.

2026-07-14 — Valarian: \$50M Series A led by NEA — its first European defense investment — with Lightbank, XTX Markets, Sequel and LitVC, plus angels Gokul Rajaram and Nikesh Arora. The London-based company (founders Max Buchan and ex-Palantir executive / U.S. Army officer Josh McLaughlin) sells a workload-level sovereignty layer for the sensitive compute that hosts AI systems and mission-critical apps, split into Valarian Enterprise and Valarian Defense editions. Total funding is now ~\$70M. It is a clean read on the emerging sovereign-AI-infrastructure theme: sovereignty enforced at the workload layer, sold to governments and regulated enterprises wary of U.S.-hyperscaler dependence.

2026-07-09 — QIZ Security: \$17M seed co-led by Bessemer and Merlin Ventures (with Evolution Equity, Qbeat Ventures, Singtel Innov8 and Qino Cyber Capital) for cryptographic posture management and post-quantum readiness. Founded May 2025 by Ben Volkow — whose Otonomo IPO'd in 2021 and whose Traffix sold to F5 — with Itan Barmes and Lenny Ridel. The round was logged in the 7/10 run's Funding Rounds tab; this run adds the Companies row (#204). It remains marginal-AI — the AI content is risk modeling — but is tracked for the crypto-posture theme.

Macro context: the week's largest data-security figure was a correction rather than a new round. Cyera's June Series G is confirmed at \$600M at a \$12B valuation, led by Evolution Equity Partners with Cyberstarts and Temasek — roughly doubling the \$6B mark set in mid-2025 and cementing AI-era data security as the highest-valued category in the tracker.

Exits & M&A

One catch-up deal was logged. Cisco announced on 2026-06-18 its intent to acquire WideField Security, a California identity-threat-detection (ITDR) startup, to strengthen Splunk's Agentic SOC — Cisco's third security acquisition of 2026 after Astrix and Galileo. Terms were undisclosed. The deal reinforces two running themes: identity as the center of gravity for AI-era security, and platform vendors buying ITDR / agentic capability rather than building it. No in-window AI-security exits, IPOs or shutdowns surfaced.

New Entrants

Four companies were added (#201–204). Oak (#201) and Valarian (#203) are covered under Funding Activity. First Recon AI (#202) launched its “AI Security Runtime” to general availability on 2026-07-08 — an inline enforcement layer that sits between users, apps and any model, inspecting every interaction (human-to-model, agent-to-tool, and agent-to-agent) and applying policy before data reaches the model, with each decision recorded as audit-ready evidence. Founded and led by Kentaro Kawamori, it exited stealth already partnered with pan-European MSP Conscia Group; funding was not disclosed. QIZ Security (#204) is covered above.

Tracked-Startup Updates

Cyera (#47) was refreshed in place: valuation raised to \$12B, total funding to ~\$1.9B, stage moved to Growth (Series G), and the June round reconciled from a previously mis-logged \$300M to the correct \$600M (2026-06-10). This was the week's one material data-quality fix — the company row had not been touched since the 2025 \$6B round despite a logged, if mis-sized, 2026 round.

Separately, ten long-standing stub rows were completed with founders, HQ and funding detail: Spectrum Security (Meny Har; \$19M, TechOperators / SkinOS), Artemis (Shachar Hirshberg and Dan Shiebler; \$70M, Felicis-led), Clearly AI (Emily and Joe Choi-Greene; \$8.4M, YC S24; HQ set to Seattle), Fig Security (Gal Shafir and team; \$38M, Ten Eleven / Team8; HQ Israel), Glide Identity (Eran Haggiag; >\$25M, Crosspoint), Humanix (Keith Stewart; \$18M, Acrew / boldstart), Trent AI (Neil Lawrence and team; \$13M, LocalGlobe / Cambridge Innovation Capital; HQ corrected U.S. → London), Ray Security (Ariel Zamir and team; \$11M, Venture Guides / Ibex), Sevii (Curt Aubley and team; autonomous defense & remediation, pre-seed), and SurePath AI (Casey Bleeker and Randy Birdsall; \$5.2M seed).

Theme Signals

Sovereign AI infrastructure / workload governance is now fundable at scale: Valarian's NEA-led round enforces sovereignty at the workload layer for AI and mission-critical compute, riding sovereignty anxiety about U.S.-hyperscaler dependence and rising defense-AI budgets. It sits adjacent to national-AISI evaluation work and the “frontier cyber as national infrastructure” thread from prior weeks.

Post-quantum crypto-agility (cryptographic posture management) is coalescing into a category — QIZ frames CPM as the CSPM/DSPM-equivalent control point for the NIST-PQC migration ahead of “Q-Day.” Marginal-AI, but the inventory-and-remediation control-point pattern is the same one that made CSPM and DSPM durable.

Frontier-lab AI-misuse defense is becoming a first-class internal function. Anthropic disclosed (Axios, 7/15) roughly 32 safeguards roles targeting chemical/explosives, nuclear, financial-scam and cybercrime misuse, plus rolling AI Safety and AI Security Fellows cohorts. Labs increasingly act as both buyers and standard-setters for AI-security tooling — and, as the Virtue AI → Meta acqui-hire showed, as absorbers of AI-security talent.

AI is deflating the vulnerability-research scarcity premium. Daniel Miessler's Unsupervised Learning #536 argues that as AI makes bug-finding cheap, the old bounty scarcity premium fades and value shifts from discovery toward triage, validation and remediation — precisely where autonomous-pentest players such as XBOW, RunSybil and Hacktron are positioning. Finally, Oak's round reinforces the existing “agent-era identity consolidation” theme: the standalone identity-control-plane window is still open, but incumbents (SailPoint, 1Password, Cisco) are buying into it fast.

Failures & Pivots

No AI-security shutdowns, layoffs or pivots surfaced in-window.

New Venture Ideas

43. Cryptographic posture management for AI/ML pipelines and agent channels — a wedge beneath QIZ's general-estate CPM, targeting the crypto surface specific to AI: the keys and certificates protecting model weights and vector stores, and the TLS / MCP / A2A channels agents use to talk to tools and each other.

44. Vendor-neutral agent-to-agent (A2A) authorization broker — a standards-based trust plane issuing verifiable agent credentials and capability tokens with per-interaction authorization, sitting between the runtime-inspection layer (First Recon) and the identity layer (Oak), and neutral across MCP / A2A protocols.

45. Sovereign AI runtime-as-a-service for the regulated mid-market — a commercial analog to Valarian's defense/nation focus, packaging data, workload and agent sovereignty plus compliance-evidence for regional banks, hospital systems and state/local government without Palantir-scale integration.

People & Thought Leaders

Six founders were added, all tied to this week's entrants: Shai Morag and Tal Marom (Oak), Kentaro Kawamori (First Recon AI), Ben Volkow (QIZ Security), and Max Buchan and Josh McLaughlin (Valarian). Morag (Ermetic, Secdo) and Volkow (Otonomo, Traffix) are notable repeat security founders; McLaughlin's ex-Palantir / U.S. Army background typifies the defense-tech provenance now common in sovereign-AI startups. On the anchor list, the week's notable activity was Anthropic's safeguards hiring surge (see Theme Signals); no anchor-name new-venture launches surfaced.

Tools & Open-Source Releases

Two projects were added. AI-Infra-Guard (A.I.G), from Tencent's Zhuque Lab, is a full-stack AI red-teaming platform that stratifies testing across the agent attack surface — OpenClaw, Agent, Skills, MCP and AI-infra scans plus an LLM jailbreak harness (26+ attack operators over 16 datasets), backed by 75+ AI-component fingerprints and 1,400+ vulnerability rules — and is billed as the first framework to audit the MCP supply chain (TechTimes, 7/6). MCP-Airlock (recently renamed Trentina) is an open-source prompt-injection defense that has grown from a web-content sanitizer into a full MCP gateway, combining heuristic-plus-classifier screening and primary/quarantine-agent separation with per-consumer profiles, tool allowlists and parameter-level access controls.

Conference Highlights

No major in-scope conference fell inside the window. Black Hat USA and DEF CON 34 (early-to-mid August) are the next tentpoles; tool-release and talk scanning will resume as their programs firm up.

Newsletter Roundup

Unsupervised Learning #536 (Daniel Miessler, 7/13): the headline argument is that AI is collapsing the bug-bounty scarcity premium (captured as a theme above); the issue also flagged a dormant GitHub account reactivating to drop a LoadMaster RCE exploit kit.

AISechub (Tal Eliyahu): the AI-security digest continues to aggregate research, reports and tooling; no un-tracked company was highlighted as notable this issue.

Venture in Security (Ross Haleliuk): recent essays reinforce two theses the tracker already reflects — that CISOs now buy remediation over detection, and that Israeli cyber dominance shapes founder dynamics globally. (The Substack archive was JavaScript-gated this run; summarized from indexed essays.)

Index hygiene

Companies: 204 entries (up from 200). Funding Rounds: 120. Exits & M&A: 47. Themes: 45. Startup Ideas: 45. People & Thought Leaders: 74. AI Security Tools & Projects: 47. Sources: 246.

Weekly Update — Week of 2026-07-06

Normal 7-day weekly run (window 2026-07-03 to 2026-07-10, attempt 1). Not a quarterly deep dive (the last, 2026-05-18, is 53 days old) and not an annual sweep (the last was 2026-06-26). The July 4 holiday made this a thin week for primary AI-security rounds; the weight of the week was governments operationalizing frontier cyber models, agentic-AI attacks moving in-the-wild, and the recovery of a June identity-M&A wave that prior runs missed (five catch-up exits logged below). Ten stubs were remediated and two factual corrections made (RunSybil founders; XBOW company row synced to its Series C).

Funding Activity

2026-07-07 — Savi Security: \$7M seed led by Acrew Capital (Magnify, TTCER, Resolute) for consumer AI-scam defense; the app launched the same day. Founded by Patrick Coughlin (TruSTAR co-founder, ex-Splunk). An AI agent can join a suspicious call live and flag behavioral 'grift' tells — voice-clone 'kidnapper' ransom calls are the headline use case. \$8/month per family.

2026-07-09 — QIZ Security: \$17M seed co-led by Bessemer and Merlin Ventures for cryptographic posture management / post-quantum readiness. Adjacent rather than core (no AI product claim), so it is logged in Funding Rounds only, not Companies.

Early July — 8Layers (Madrid): \$1.1M pre-seed extension (total \$2.9M) from Criteria Venture Tech and Bankinter. ISPM + ITDR that inventories human, service-account, API and AI identities via Okta / Entra / Google Workspace / AWS IAM connectors.

Catch-up rounds recovered into the Funding Rounds tab: Armadin \$189.9M seed + Series A (Accel lead, 2026-03-10 — claimed largest seed+A in cybersecurity history); RunSybil \$40M (Khosla, 2026-03-18); Prophet Security \$11M strategic from Amex Ventures + Citi Ventures (2026-02-25, lifting total to \$41M); Realm Labs \$5M (Crosspoint, RSAC-timed). The Tolmo row was fixed in place: \$22M seed led by Accel (was n/d).

Macro: Momentum Cyber's mid-year review (7/1) counted a record 219 cyber M&A deals / \$9.1B disclosed in H1 2026 and declared AI security 'a dominant M&A reality'; the week's generalist VC roundups show capital concentrating in physical AI and agentic fintech rather than security.

Exits & M&A

In-window: Barracuda acquired Evo Security (7/7, undisclosed) — IAM/PAM built for MSPs folding into BarracudaONE, pitched as 'identity resilience for the AI era'. Marginal-AI but logged for the identity-consolidation theme.

Recovered June wave (missed by the 6/19 and 6/26 runs): SailPoint to acquire Entro Security (~\$200M reported, 6/18) — a tracked company, row updated; Accenture took a majority stake in Dragos plus all of runZero and NetRise (~\$4.175B combined EV, ~\$208M combined ARR, 6/18) — runZero is HD Moore's company; Databricks acquired Panther (6/16), its third security buy; 1Password acquired Apono (~\$250-300M reported, 6/15) to become a 'control plane for agentic AI'; Aikido acquired Root (6/30), whose AI agents auto-patch open-source vulnerabilities. No AI-security IPOs or shutdowns in-window. Macro context: Anthropic has filed a confidential S-1 (targeting ~Oct 2026); OpenAI reportedly delaying an IPO to 2027. Big-tech talent watch (Step Z): no new AI-security acqui-hires into frontier labs this week; the 7/9 Tech Times follow-up on Meta's absorption of the Virtue AI team confirms Carlos Guestrin did NOT join Meta and remains at Stanford as Virtue's Chief Scientist.

New Entrants

Savi Security (#199, Los Angeles) — first consumer-facing 'agentic call-defense' product; behavioral AI for call/message screening with live agent-in-call analysis; model stack not disclosed.

8Layers (#200, Madrid) — EU-regulatory-native identity posture that treats AI agents as first-class identities; under the hood not disclosed.

Identity resolved: SSIL (#133) is operating as Abundant Security, Inc. (ssil.ai) — Hyrum Anderson (CEO; ex-Cisco Foundation AI, ex-CTO Robust Intelligence), Joshua Saxe (CTO; ex-Meta Llama security lead), Gabe Gottlieb (COO). Framing: 'defining, measuring, and advancing security superintelligence for defenders'; products 'coming soon'; listed by F4 Fund; funding still undisclosed. Saxe speaks at USENIX Security '26 on building agents for cyber defense.

Noted but out of core scope: Arkenstone Defense (\$35M seed, J2 Ventures — US-government security compliance for commercial tech).

Tracked-Startup Updates

XBOW — company row synced to reality: Series C \$120M at \$1B+ valuation (2026-03-18, DFJ Growth + Northzone; total ~\$237M) plus the \$35M strategic extension (NVIDIA, Accenture, Samsung, SentinelOne); Ramin Sayar joined the board.

Lakera — Check Point acquisition (~\$300M) confirmed CLOSED ~Q4 2025; Zurich HQ is now Check Point's global AI-security R&D center. Torq — Series D \$140M at \$1.2B (Jan 2026, total \$332M, ~300% 2025 revenue growth, IPO ambitions); a PitchBook record shows Torq acquired Jit (AI AppSec agents) on 5/19 — not yet independently confirmed. Entro Security — stage moved to Acquired (announced, SailPoint).

RunSybil — founder correction: Ari Herbert-Voss (CEO; OpenAI's first security research hire) and Vlad Ionescu (CTO, ex-Meta red team), not Dane Stuckey as previously recorded (Stuckey is OpenAI's CISO). Customers include Cursor, Notion, Turbopuffer, Baseten and Thinking Machines Lab. Noma Security — launched Agentic Access Control for AI agents and MCP servers (June); its \$100M 2025 round was Series B (total \$132M). Prophet Security — total funding now \$41M after Amex/Citi strategics. Horizon3.ai — Series D was \$100M led by NEA, closed 2025-06 (total \$178.5M; 3,000+ orgs, 150K+ autonomous pentests).

Dropzone AI — expanding from AI SOC analyst to a full agentic SOC (AI Threat Hunter + AI Threat Intel Analyst, Q2 2026); Edward Wu named to Fast Company's 'AI 20 for 2026'. Simbian — AI Threat Hunt Agent + Microsoft Sentinel data-lake integration; claims the first reasoning-based platform unifying offensive and defensive security. Zenity — Carahsoft public-sector distribution (NASA SEWP V, TIPS) and two categories of Gartner's 2026 Agentic-AI Hype Cycle.

Stub remediation (10): Realm Labs (Sunnyvale; Saurabh Shintre + Akash Mukherjee; \$5M Crosspoint-led; white-box LLM 'internal observability' inspecting attention patterns and token probabilities; Anthropic an early customer); Charm Security (NY/Team8; Roy Zur + Avichai Ben David; \$8M); Crash Override (NY; John Viega + Mark Curphey; \$28M GV/SYN seed); Gecko Security (\$500K, YC F24); AISLE (funding still undisclosed; angels include Jeff Dean and Olivier Pomel; found 12 new OpenSSL vulns in Jan); Abundant/SSIL (above); Tolmo (SF; ex-Sqreen/Datadog founding team; \$22M Accel seed; agents over a live production knowledge graph); Twine Security (Tel Aviv; Benny Porat, Claroty co-founder/CTO); Portkey (\$18M total; PANW acquisition intent stands); Armadin (SF; \$189.9M; co-founders Travis Lanham, Evan Peña, David Slater).

Theme Signals

Frontier cyber models as national infrastructure. OpenAI shipped GPT-5.6 'Sol' (7/9) as its 'most capable cybersecurity AI' after a 12-day White House safety review; Reuters reports CISA is using Anthropic's Mythos to scan federal code for flaws (NSA reportedly a user too); the UK launched 'Cyber Shield' — national-scale agentic defense with red/blue agent teams — alongside a 60-signatory Cyber Resilience Pledge (7/7).

Agentic-AI attacks moved in-the-wild — the week's dominant research theme. 'GitLost' (Noma Labs, 7/8): unauthenticated prompt-injection exfiltration via GitHub Agentic Workflows; actor 'JadePuffer' ran an LLM mid-intrusion for agentic ransomware via Langflow CVE-2025-3248; Zscaler showed agents tricked into unauthorized crypto payments; 'DuneSlide' zero-click Cursor RCEs (CVSS 9.8). DEF CON 34's AI Village added an 'Adversarial Attacks Against Agents' poster track.

'Securing AI' is now a Gartner budget line — and the only accelerating one. The 2Q26 forecast breaks out securing-AI spending for the first time: \$21.9B of \$154.4B in new security spending through 2030, growth rising every year (16.3% to 20.1%), passing endpoint protection by 2029 (via Software Strategies Blog, 7/6).

Agent-era identity consolidation: four identity/access deals in four weeks (SailPoint-Entro, 1Password-Apono, Cisco-Astrix closed, Barracuda-Evo) plus Noma's Agentic Access Control launch. The standalone NHI/agent-IAM window is narrowing; remaining independents (Token, Oasis, Clutch, Keycard, Twine) look like scarce assets.

Failures & Pivots

Snyk — layoffs plus an AI pivot (announced ~6/24, surfaced in SecurityWeek's 7/3 roundup): R&D realigned around four AI-security areas (AI-written code, autonomous agents in production, vulnerability chaining, always-on adversaries). Headcount figures conflict: ~198 (14%) per Calcalist vs 90 per Globes. Follows CEO Peter McKay's February departure.

New Venture Ideas

40. Untrusted-content provenance firewall for agent workflows

GitLost, AutoJack and SearchLeak all exploit the same primitive: untrusted content (public issues, webpages, search results) entering agent context with ambient authority. Wedge: an inline provenance gateway that labels and quarantines untrusted spans and enforces per-tool trust-handoff policies before agent execution. Buyers: platform teams running GitHub Agentic Workflows and Copilot/Cursor fleets. Moat: provenance taxonomy plus an incident-derived detection-content flywheel. Complements gateway-visibility plays (including Helmet) — this is the content-lineage layer, not the guardrail layer.

41. Sovereign frontier-cyber evaluation and attestation broker

Governments now gate frontier cyber models (the White House's 12-day GPT-5.6 review; UK Cyber Shield; CISA's operational use of Mythos) but lack independent, repeatable evaluation. Build the ExploitBench-style certification harness and attestation chain, sold to national programs, frontier labs seeking pre-cleared deployment, and defense primes. Moat: accredited methodology plus clearances — neutral-party positioning the labs cannot replicate themselves.

42. VIP 'digital bodyguard' for executives and their families

Savi proves consumer demand for AI-scam call defense at \$8/month; enterprises face exec-targeted deepfake ransom and voice-clone attacks, but deepfake vendors (Adaptive, Reality Defender) focus on workforce training and detection, not executives' households. Wedge: B2B2C executive protection — live call-join agents and family device coverage sold through CISOs and family offices at 100-1000x consumer ACV. Risk: Savi moves up-market first.

People & Thought Leaders

New entries (5): Hyrum Anderson (Abundant Security CEO); Patrick Coughlin (Savi CEO); Ari Herbert-Voss (RunSybil CEO); Benny Porat (Twine CEO, Claroty co-founder); Pliny the Liberator / elder-plinius (pseudonymous AI red-teamer; released T3MP3ST).

Updates: HD Moore (runZero to Accenture); Matt Suiche (now founder of OnDB — data infrastructure for the agentic economy); Vijay Bolina (CISO at a stealth physical-AI frontier lab, Prometheus Inc. per Crunchbase; BH USA '26 AI Summit bio); Edward Wu (Fast Company AI 20); Itamar Golan (reportedly building a new venture, unconfirmed); Carlos Guestrin (stays Stanford/Virtue); John McIntosh (RECon 2026 workshop on agentic reverse engineering); Jason Clinton (correction: Deputy CISO at Anthropic since 9/30/25 under CISO Vitaly Gudunets).

Tools & Open-Source Releases

In-window: T3MP3ST (Pliny/elder-plinius, ~7/6) — multi-agent offensive-security meta-harness; Nika (PhonePe, 7/1) — cross-file taint analysis for Java microservices.

June wave now tracked (13): AIMap (exposed Ollama/MCP endpoint discovery and attack-testing); AgentGG (agentic SAST); Agent Beacon (Asymptote Labs — coding-agent telemetry); Agent Threat Rules (open detection-rule format for agent threats); OWASP Agent Memory Guard (screens agent memory read/writes); Praxen (verifies agent behavior against declared policy); DarkMoon (autonomous pentest); Microsoft Clarity + RAMPART (AI Red Team agent design review + continuous testing); OpenHack (Hadrian, agent-driven vuln research); AgentStalker (agent-vuln benchmark); SkillsGuard (malicious skill-package scanner); agent-jackstop (Tenet Security — hardening vs agentjacking); mcp-trust-plane (Abluva — MCP guardrail plane).

Conference Highlights

Black Hat USA 2026 (Aug 1-6) briefings are live, and the AI-offensive lineup is strong: Novee on breaking AI agents in official workflows (trust-handoff failures reproducible across Anthropic, Google and OpenAI agent stacks); Rein Security on hacking a top-3 US retailer's AI shopping assistant; Rubrik Zero Labs on a Copilot sandbox-to-host escape ('ChatMate'); and NVIDIA showing a fine-tuned 30B open-source model achieving 56% exploit success at 70-125x lower cost than frontier models. AI Summit Aug 4.

DEF CON 34 (Aug 6-9): AI Village confirmed with demos, a live-system CTF, and the new adversarial-agents poster track. RECon 2026 (June, catch-up): John McIntosh ran 'Agentic

Reverse Engineering: Building Custom AI Skills with Coding Agents'. USENIX Security '26: Joshua Saxe presenting 'Lessons Learned in Building Agents for Cyber Defense'.

Newsletter Roundup

Unsupervised Learning #535 (7/6): leads with Pliny's T3MP3ST offensive harness and Miessler's LifeOS 6.x (with new 'Conduit' local context system). AISechub June edition (6/28): dense agent/MCP supply-chain issue — SearchLeak (M365 Copilot one-click exfil), AutoJack (webpage-to-RCE on agent hosts), Miasma Worm npm compromise, six critical AI-infra CVEs (mcp-pinot and n8n MCP Browser both CVSS 10.0), plus the OSS tool wave captured above. Venture in Security (6/23): essay on invisible security leaders — no company callouts. Software Strategies Blog (7/6): the Gartner securing-AI breakout analyzed under Theme Signals.

Helmet Watch

No news in-window (no funding, hires, or press). One soft signal: helmet.sh positioning now reads 'visibility and control over agents, MCP servers, skills, plugins, and gateways' — the 'skills, plugins, and gateways' language suggests scope has broadened beyond MCP-only messaging (change undated). Blog unchanged since ~January; the \$9M seed baseline (SYN Ventures + WhiteRabbit VC, Dec 2025) stands.

Index hygiene

Companies: 200 (+2). Funding Rounds: 118 (+7 appended, 1 fixed in place). Exits & M&A: 46 (+6). Themes: 41 (+4). Startup Ideas: 42 (+3). People & Thought Leaders: 68 (+5 new, 8 notes updated). AI Security Tools & Projects: 45 (+15). Sources: 218 (+25). Stubs remediated: 10. Corrections: RunSybil founders; XBOW row sync; Tolmo round.

Weekly Update — Week of 2026-07-03

Normal 7-day weekly run (window 2026-06-26 to 2026-07-03, attempt 1). A quiet week immediately after the 2026-06-26 annual-sweep combo, so deal flow was thin and the emphasis this run is on catch-up entrants that predated the sweep, in-window product launches, and stub remediation. This was not a quarterly deep dive (the last, 2026-05-18, is 46 days old) and not an annual sweep (the last was 2026-06-26).

Funding Activity

No brand-new in-window AI+security funding rounds surfaced — expected for the week following a large sweep.

2026-03-10 — Quantro Security: \$2.5M seed (Gradient, Google's early-stage AI fund) at a ~\$25M valuation for autonomous vulnerability management. Catch-up (pre-sweep, previously untracked).

2026-05-15 — Hacktron: \$2.9M pre-seed for AI-driven continuous penetration testing that tests every code change like an attacker. Catch-up (previously untracked).

Exits & M&A

No new in-window AI+security exits, M&A or acqui-hires.

Big-tech / frontier-lab talent watch (Step Z): the week's major lab moves were general-AI, not AI-security. Anthropic's ~\$400M all-stock acqui-hire of Coefficient Bio is computational biology, and OpenAI's hire of policy adviser Dean Ball (Head of Strategic Futures, starting 7/6) is AI-policy — neither is a security-startup acqui-hire, so nothing was logged in Exits & M&A.

New Entrants

Quantro Security (New York; founded 2025 by CrowdStrike / Tenable / Qualys veterans) — emerged from stealth 3/10/26 with \$2.5M seed from Gradient at a ~\$25M valuation. Its VM.Analyst agent ingests CMDB, cloud, firewall and VM data, then reasons over an org's real compensating controls to rank which vulnerabilities are genuinely exploitable, with natural-language task delegation. Catch-up.

Hacktron (San Francisco; competitive-hacker founders with Cambridge ties) — \$2.9M pre-seed; AI that tests every code change like a real attacker, with disclosed testing for Perplexity and Supabase. Catch-up.

Netzilo (Campbell, CA; CEO Egemen Tas, ex-Comodo / Xcitium CTO) — in-window: on 7/1 expanded its 'AI control plane for the agentic workforce,' adding runtime governance (AIDR) across Amazon Bedrock AgentCore, Microsoft Foundry, Copilot Studio, CrewAI, LangGraph and Vertex AI, and open-sourced its AIDR Sigma detection rules.

Tracked-Startup Updates

Exaforce (#117): refreshed — added USP and under-the-hood detail (real-time security knowledge graph + autonomous 'Exabot' agents orchestrating LLM reasoning; platform + MDR). The \$125M Series B at \$725M val / \$200M total was already on file. Flagged the pre-existing duplicate row #12 (old Seed stub) for later reconciliation.

Surf AI (#134): stub remediated — Tel Aviv (founded 9/2024); founders Yair Grindlinger (CEO), Elad Horn, Roie Cohen Duwek, Avner Gideoni, Brenton Gumucio; \$57M led by Accel (Cyberstarts, Boldstart); context-graph agentic security ops.

Gambit Security (#137): stub remediated — founders Alon Gromakov (CEO), Sa'ar Elias (CPO), May Kogan (CTO), Unit 8200 vets ex-Sentra; \$61M (\$5M seed + \$56M Series A; Spark Capital, Kleiner Perkins, Cyberstarts); AI-native cyber-resilience.

Onyx Security (#105): stub remediated — founders Maxim Bar Kogan and Gil Elbaz; \$40M (\$5M seed 2024 + \$35M Series A led by Conviction; Cyberstarts); secure AI control plane for agent governance; Tel Aviv + NY.

Novee (#107): stub remediated — founders Ido Geffen (CEO), Gon Chalamish (CPO), Omer Ninburg (CTO), IDF elite cyber vets; Tel Aviv; \$51.5M (\$8.5M seed + \$33M Series A + \$10M venture debt; YL Ventures, Canaan, Oren Zeev); 'proprietary AI hacker' for continuous offensive testing.

Compliance (#155): stub remediated — founder/CEO Richa Kaul; \$20M Series A led by GV (Google Ventures); ~\$28M total; agentic AI GRC deploying 30+ specialized compliance agents (HIPAA/ISO/NIST).

Theme Signals

AIDR (AI Detection & Response) is emerging as a named runtime-agent-security pattern, now with an open detection-rule ecosystem: Netzilo open-sourced Sigma-style AIDR rules (github.com/netzilo/aidr-sigma) modeling multi-stage agent behavior (tool calls, file reads, skill acquisition) to catch prompt injection, tool poisoning and staged exfiltration — the 'Sigma for agents' idea moving from concept to code.

Shadow-AI discovery / governance is productizing down to the endpoint and browser. In-window GA launches from incumbents: Jamf AI Governance (discover/enforce/report on AI tools on Mac) and iboss's free AI Security Platform (real-time visibility into ChatGPT, Copilot, Gemini, Claude, Perplexity, Cursor). The category is commoditizing fast, squeezing pure-play shadow-AI startups.

Autonomous offensive security keeps drawing marquee capital and founders (A Security, Novee, Hacktron, XBOW, Dreadnode) — 'AI hacker' continuous-pentest positioning is now a crowded lane.

Failures & Pivots

No AI+security startup shutdowns, layoffs or pivots identified in-window.

New Venture Ideas

Idea #38 — Vendor-neutral AIDR rule exchange + managed detection. Netzilo open-sourcing AIDR Sigma rules echoes the early Snort/Sigma moment; build the neutral registry + curation + managed-detection layer for agent-runtime behavior across every harness (Bedrock AgentCore, Foundry, LangGraph, CrewAI, Vertex), monetizing content and response rather than a single agent.

Idea #39 — Cross-platform shadow-AI governance for the regulated mid-market. As Jamf (Mac) and iboss (network) commoditize discovery, a unified control plane spanning endpoint + browser + network + SaaS with audit-ready evidence mapped to HIPAA/ISO/NIST could win the compliance-driven buyer the point tools fragment.

People & Thought Leaders

New entry: Egemen Tas (Netzilo, CEO/founder; ex-Comodo / Xcitium CTO) — building the AIDR / agentic-workforce control plane.

People → company tripwire (Step P): no further unlisted anchor-name ventures surfaced in-window; the week's frontier-lab moves (Coefficient Bio → Anthropic; Dean Ball → OpenAI) are biotech and policy, not AI-security founders.

Tools & Open-Source Releases

Netzilo AIDR Sigma rules (github.com/netzilo/aidr-sigma, ~7/1/26): community-released AI Detection & Response rules encoding multi-stage agent-behavior signatures (prompt injection, indirect injection, tool poisoning, capability hijacking, privilege escalation, staged exfiltration) for Bedrock AgentCore, Foundry, Copilot Studio, CrewAI, LangGraph and Vertex.

No in-window RE / vuln-research tool drops — RECon (6/19-21) was covered last run, and DEF CON 34 and Black Hat USA land in early August.

Conference Highlights

No major security conference fell in the 6/26-7/3 window. Watch list: Black Hat USA and DEF CON 34 (Las Vegas, early-mid August) and Hexacon 2026 (Paris/Oct, opening with an 'AI for vulnerability research' keynote).

Newsletter Roundup

Venture in Security (Ross Haleliuk): recent programming centered on 'build vs. buy' for AI security stacks and AI-powered credential security; post-RSAC, Haleliuk stresses practitioners' real worries are shadow AI and attackers using AI for recon at scale against identity / vuln-management gaps — no new funded-company names to capture.

Help Net Security 'new infosec products' (7/3): featured Netzilo (AIDR runtime governance), Jamf AI Governance and iboss's free AI Security Platform — all captured above.

Unsupervised Learning / AISecHub / Software Strategies Blog: no new in-window issues surfacing untracked funded companies since the 6-26 run.

Helmet Watch

No Helmet Security (helmet.sh) news in-window — no new funding, hires, product launches or press. The latest remains the December 2025 \$9M seed (SYN Ventures, WhiteRabbit) for MCP-connection security.

Index hygiene

Companies tab now holds 198 entries (up from 195): added Quantro Security (#196), Hacktron (#197) and Netzilo (#198). Five stubs remediated in place (Surf AI, Gambit Security, Onyx Security, Novee, Compliance) and Exaforce #117 refreshed. Funding Rounds 111 (+2), Exits & M&A unchanged at 40, Themes 37 (+2), Startup Ideas 39 (+2: #38, #39), People 63 (+1: Egemen Tas), AI Security Tools & Projects 30 (+1: Netzilo AIDR Sigma), Sources 193 (+13). Pre-existing Exaforce #12/#117 duplicate flagged for reconciliation.

Weekly Update — Week of 2026-06-26 (Annual Historical Sweep)

Combined run: the normal 7-day weekly window (2026-06-19 to 2026-06-26, attempt 1) plus the first Annual Historical Sweep - no prior ANNUAL_SWEEP had ever been logged. This was not a quarterly deep dive (the last deep dive, 2026-05-18, is only 39 days old). The sweep

deliberately looked backward to catch solid 2024-2025-funded AI-security companies, several with marquee research/founder pedigrees, that predate this tracker's April 2026 start - the same blind spot that hid Virtue AI until Meta acqui-hired it. Net result: 21 companies added (2 in-window, 5 recent-month catch-ups, 14 historical/geographic backfill), 5 existing rows refreshed, 1 missed exit recovered (Koi to Palo Alto Networks), 22 funding rounds, 12 open-source tools, 10 people, 5 themes and 3 ideas. Connectivity note: the curl preflight returned HTTP 403 (raw egress is blocked in this environment), but the research tools route separately and were verified live, so the run proceeded.

Funding Activity

2026-06-24 - Runlayer: \$30M Series A (Felicis, with Khosla Ventures) for an agent-governance 'control layer' sitting between employees / AI agents and enterprise tools, data and policies; total \$42M. Customers include Instacart, Gusto, Decagon, Opendoor and dbt Labs.

2026-06-17 - Tenet Security: \$6M seed (The Westly Group, MizMaa Ventures) to prevent attacks on enterprise AI agents using pre-execution 'Agent-side Simulation.'

2026-06-17 - NeuralTrust: EUR 17.2M (~\$20M) seed (Alstin Capital), billed as the largest EU cybersecurity seed to date, to secure and govern enterprise AI agents (Barcelona).

2026-06-04 - Opal Security: \$23M (Greylock, Battery Ventures; total \$59M) for AI-native access governance spanning human, non-human and agentic identities; launched 'Agents in Opal' at Identiverse 2026.

2026-06-02 - ZeroDrift: \$10M seed for a 'compliance firewall' that checks AI-generated email, chat, voice and video against regulatory rules before enterprise systems send them.

Backfill rounds added to the tracker (predating the sweep window): Tenex.AI \$250M Series B (3/31, total \$277M), Cylake \$45M seed (Nir Zuk), Irregular \$80M (~\$450M val), depthfirst ~\$120M across two 2026 rounds, Dreadnode \$14M, Keycard \$38M, Aurascape \$50M, Singulr AI \$10M, plus General Analysis, NanoCo, Secludy, Symbiotic Security, Jericho Security, Prime Security, Operant AI, Onit Security and ZAST.AI. About 22 rounds in all.

Exits & M&A

No new in-window AI+security exits this week. The Virtue AI to Meta Superintelligence Labs acqui-hire (6/25) was already logged in last week's run.

Backfill recovery: Palo Alto Networks completed its ~\$400M acquisition of Koi (4/14/26), forming 'Agentic Endpoint Security' (AES) now integrated into Prisma AIRS and Cortex XDR. Koi (Israeli, founded 2024 by Unit 8200 alumni, \$48M raised) had been a tracker stub; the row is now remediated and marked acquired. The deal had been missed in prior weekly runs.

New Entrants

In-window: Runlayer (agent governance / MCP control layer) and Tenet Security (agent runtime security via 'Agent-side Simulation').

Recent-month catch-ups now tracked: Opal Security (AI-native access governance), ZeroDrift (AI-communications compliance firewall), General Analysis (agentic AI red-teaming; founders ex-NVIDIA / Cohere / DeepMind), NanoCo (NanoClaw secure local-first agent assistant), Secludy (privacy-preserving synthetic data for GenAI).

Annual-sweep backfill - frontier / research-founded labs and pre-tracker companies: Irregular (ex-Pattern Labs; frontier AI security lab; SOLVE framework; Dan Lahav / Omer Nevo); Dreadnode (autonomous AI red-teaming; Will Pearce, ex-Microsoft/NVIDIA AI red teams); depthfirst (AI app/code security; ex-DeepMind co-founder); Aurascape (shadow-AI data protection; ex-Zscaler); Prime Security (AI security-by-design); Operant AI (runtime cloud + AI

security); Jericho Security (GenAI / deepfake phishing defense); Singulr AI (enterprise AI governance).

Geographic backfill: Keycard (Halifax, Canada - agent identity / IAM; a16z, boldstart, Acrew); Symbiotic Security (NY / France - security for AI-generated code); Fencio (India - runtime agent guardrails); Cylake (Israel - Nir Zuk's AI-native, sovereign on-prem SecOps, \$45M from Greylock).

Tracked-Startup Updates

Koi (#96): stub remediated and marked acquired by Palo Alto Networks (~\$400M, 4/14/26) -> Palo Alto 'Agentic Endpoint Security'; founders Amit Assaraf, Idan Dardikman and Itay Kruk added.

Tenex.AI (#110): refreshed - \$250M Series B closed 3/31/26 (total \$277M); founders (Eric Foster / ex-Cyderes, Venkata Koppaka / ex-Google SecOps) and Sarasota HQ filled in.

NeuralTrust (#148): stub remediated and EUR 17.2M seed (6/17) recorded; founders and Barcelona HQ added.

Qevlar AI (#151): stub remediated - Paris HQ, founders Ahmed Achchak and Hamza Sayah, \$30M (Partech / Forgepoint / EQT Ventures).

Token Security (#97): founders (Itamar Apelblat, Ido Shlomo) and \$27M total filled in; flagged as a likely duplicate of row #58 for later reconciliation.

Theme Signals

Frontier AI security labs are now a distinct category. Irregular, Virtue AI, General Analysis and Dreadnode show specialist labs that red-team and harden frontier models in partnership with the model builders - and are prime acqui-hire targets (Virtue -> Meta).

Agent identity & authorization is consolidating into a control plane. Keycard, Opal, NewCore, Token Security and Astrix are racing to own verifiable agent identity and scoped, delegated authorization.

Runtime agent interception / egress firewalls are proliferating. Tenet's pre-execution simulation, plus open-source egress tools (Pipelock, Gen Digital's Sage, APort), point to the agent tool-call as the new enforcement point.

Agentic reverse engineering went mainstream at RECon 2026. The 6/19-21 program was themed around agent-driven RE, with a wave of headless MCP servers for IDA / Ghidra / Binary Ninja and the first agentic-RE benchmark (AgentRE-Bench).

Sovereign / on-prem AI security is drawing marquee founders. Nir Zuk's Cylake (\$45M, no public cloud) and EU sovereign-inference plays signal data-sovereignty as a fundable wedge.

Failures & Pivots

No AI+security startup shutdowns, layoffs or pivots were identified in-window. Negative signal of note: OffensiveCon 2026 (Berlin) ran its full program with zero AI-security talks - a reminder that the deep VR / exploitation community remains deliberately AI-skeptical even as RECon leans in.

New Venture Ideas

Idea #35 - Portable 'agent dry-run' simulation standard. A vendor-neutral pre-execution simulation / intervention layer (the Tenet capability, but cross-framework and policy-driven); MCP standardization makes a shared 'likely next action' corpus feasible.

Idea #36 - MCP supply-chain trust registry & attestation. The MCP-server ecosystem has scanners (Cisco, GitHub) but no signing / attestation / reputation layer; build the trust registry and become the npm-audit of agent tools.

Idea #37 - Hosted agentic vuln-research pipeline for defenders. Productize the fragmented RECon-era OSS (headless MCP + AgentRE-Bench + hallucination guards) into an audited, hosted cyber-reasoning service for product-security teams.

People & Thought Leaders

New entries (10): Nir Zuk and Udi Shamir (Cylake; founders of Palo Alto Networks and SentinelOne respectively); Dan Lahav and Omer Nevo (Irregular); Will Pearce and Nick Landers (Dreadnode); Barak Sternberg and Nevo Poran (Tenet; ex-Cisco AI Defense); Moinul Khan (Aurascape; ex-Zscaler); Ian Livingstone (Keycard; ex-Manifold). Continues the practice of adding researcher / operator founders before they are widely indexed - the explicit Virtue-AI lesson.

People -> company tripwire: no further unlisted anchor-name ventures surfaced this week beyond those captured above.

Tools & Open-Source Releases

RECon 2026 (6/19-21, in-window): IDASQL (Elias Bachaalany - a SQL interface to IDA Pro, agent-agnostic); Tim Blazytko's headless MCP suite (binary-ninja-headless-mcp, ghidra-headless-mcp, and an agentic-malware-analysis Claude skill); AgentRE-Bench (an agentic-RE benchmark with a hallucination penalty).

Vulnerability discovery: OpenAnt (Knostic - two-stage LLM detect-then-adversarially-verify) and FuzzingBrain (a DARPA AlxCC finalist that claims 62 zero-days).

Agent / MCP defense: Cisco AI Defense MCP Scanner; Pipelock (agent egress firewall); Gen Digital's Sage (runtime Agent Detection & Response); from [un]prompted 2026, SYARA / SuperYARA (Palo Alto Unit 42) and Sondera's Cedar-policy coding-agent hooks; and AIOpsDoom / AIOpsShield (USENIX Security 2026). Twelve tools were added to the tracker.

Conference Highlights

RECon 2026 (Montreal, 6/19-21) - the richest in-window source; the entire program leaned into agentic reverse engineering, including talks on agentic RE of Claude's security architecture, deobfuscation in the age of agentic RE, and 'SELECT * FROM binary.'

OffensiveCon 2026 (Berlin) - notably ran with zero AI-security talks (a negative signal worth tracking).

Upcoming watch: Hexacon 2026 (Oct) opens with an 'AI for vulnerability research' keynote from Bruce Dang and Thai Duong (Calif.io), who used Anthropic's gated 'Mythos' model to build the first public macOS-kernel memory-corruption exploit on Apple's M5; Black Hat USA and DEF CON 34 (Aug) both have dedicated AI tracks; and USENIX Security 2026 has a cluster of LLM-agent security papers.

Newsletter Roundup

Venture in Security (Ross Haleliuk, 6/23): an essay arguing the most visible security voices on LinkedIn / podcasts are not representative of the practitioners actually running large security teams. No new companies flagged.

Unsupervised Learning (Daniel Miessler, No. 533, 6/20): theme of 'the harness matters more than the model,' with a strong personal nod to Scanner.dev (security data lake; Sequoia Series A) and CyCognito (AI pentesting) - both adjacent / established, so not added as new entrants. AISecHub (Tal Eliyahu, May issue): a research / tools digest focused on prompt-injection-to-RCE in AI coding agents and MCP security; no funded-company highlights this cycle.

Software Strategies Blog: the most recent cybersecurity post is April 8 (the blog is not weekly); it surfaced AI-native vuln names Onit Security and ZAST.AI (both added) and Jericho Security (added).

Helmet Watch

No Helmet Security (helmet.sh) news in-window - no new funding, hires, product launches or press. The latest remains the December 2025 \$9M seed (SYN Ventures, WhiteRabbit Ventures); its MCP-layer agent-communications security positioning is unchanged.

Index hygiene

Companies tab now holds 195 entries (up from 174). Funding Rounds 109, Exits & M&A 40, Themes 35, Startup Ideas 37, People & Thought Leaders 62, AI Security Tools & Projects 29, Sources 180. Five new themes and three new ideas were added. Next annual sweep is due ~2027-06; next quarterly deep dive ~2026-08-16.

Weekly Update — Week of 2026-06-19

Normal weekly run (7-day window, 2026-06-12 to 2026-06-19; attempt 1). A heavy week for AI-security capital: three stealth emergences (Ent, NewCore, Magnitude) plus two cyber mega-rounds (Dream, Twenty Technologies), and Chainguard's launch of the Athena coalition — an industry-run clearinghouse to fix the vulnerabilities frontier AI now finds at scale. Tracker totals after this run: 170 companies, 84 funding rounds, 38 exits, 29 themes, 34 ideas, 42 people.

Funding Activity

2026-06-16 — Ent (Ent.AI): emerged from stealth with a \$100M seed led by Decibel (with Sequoia, Crosspoint Capital, Craft Ventures, Shield Capital, Felicis and In-Q-Tel) — among the largest cyber seed rounds on record. Founders Elias Manousos and Brandon Dixon built RiskIQ (acquired by Microsoft) and worked on Microsoft Security Copilot.

2026-06-15 — NewCore: exited stealth at a \$300M valuation on \$66M from Cyberstarts, Index Ventures and Evolution Equity Partners. Founded by Zohar Alon (Dome9 → Check Point), it governs human, machine and AI-agent identities under one architecture; debuted at Identiverse.

2026-06-16 — Magnitude: launched from stealth with a \$10M seed led by Ballistic Ventures to deploy an 'autonomous AI workforce' for third-party risk management. Team drawn from Amazon, Abnormal AI, Proofpoint and Pandora.

2026-06-18 — Dream: the Israeli sovereign-AI cyber-defense firm (co-founders Shalev Hulio, ex-NSO Group, and former Austrian chancellor Sebastian Kurz) raised \$260M at a \$3B valuation (Bicycle Capital, Group 11, Antler, Bain Capital Ventures, Tru Arrow), on roughly \$300M of 2025 sales into governments and critical infrastructure.

2026-06-17 — Twenty Technologies: raised a \$100M Series B at a \$1B valuation (Accel, Caffeinated Capital, Friends & Family Capital, Point72 Ventures) to automate offensive cyber operations for the US military and intelligence community.

Context: TechStartups' June 18 tape framed the week as 'capital chasing scarcity, not novelty' — sovereign cyber defense and AI-governed enterprise control points absorbing the largest checks while generic AI positioning no longer clears the bar.

Exits & M&A

No new in-window AI+security exits this week.

Already tracked, resurfaced in this week's coverage: Akamai's ~\$205M acquisition of LayerX (secure enterprise browser / AI-usage control; close expected Q3 2026) and Check Point's earlier purchase of agent-identity startup Cyata.

New Entrants

Ent (ent.ai) — intent-aware endpoint security: a lightweight endpoint agent adds real-time AI reasoning that interprets both human and AI-agent intent and intervenes before risky actions finalize, extending EDR/SIEM/SOAR/IAM rather than replacing them.

NewCore — security-first identity for the agentic workforce: treats AI agents as first-class identities (lifecycle, trust score, revocation path), with 'Secure Split Key' signing, phishing-resistant auth and hardware-bound credentials; built for environments where agent identities outnumber humans by two orders of magnitude.

Magnitude — autonomous third-party risk: AI risk agents continuously gather and validate evidence, decide risk and drive remediation across third- and Nth-party ecosystems, increasingly governing partners' AI agents too.

Theme Signals

Private AI-vuln clearinghouses go live. Chainguard launched Athena (6/16), a 30+ member coalition (BNY, JPMorganChase, Cisco, Cloudflare, Docker, Kyndryl, LTIMindtree, PwC) that pools vulnerabilities found by frontier AI (Anthropic's Glasswing/Mythos, OpenAI's Daybreak), privately patches open source and pushes pre-disclosure mitigations — already 20,000+ findings and 2,000+ patches across 500 projects. It explicitly models the 'AI cybersecurity clearinghouse' the June 2 Executive Order asked the US government to build.

Intent-aware endpoint is the new agentic control point. Ent's \$100M seed reframes endpoint security from post-hoc detection to real-time intent evaluation at the human/agent action boundary.

Sovereign and offensive AI-cyber draw mega-rounds. Dream (\$3B) and Twenty Technologies (\$1B) show capital concentrating on nation-scale cyber where government demand is durable and board-defensible.

Autonomous third-party risk (agentic GRC) is forming as a category. Magnitude joins Lema AI in replacing questionnaire-bound TPRM with autonomous evidence-gathering, risk decisions and remediation.

Agent-identity valuations keep climbing. NewCore's \$300M de-stealth valuation — on top of Offroad, Willow and Geordie — shows seed-stage agent-identity now priced like later-stage rounds.

Data-to-model governance heats up. Infosecurity reported (6/17) that sensitive enterprise data uploads to AI models doubled year-on-year, reinforcing the DSPM-for-AI thesis.

Failures & Pivots

No AI+security startup shutdowns or pivots identified in-window. Macro note only: large-company security-team layoffs were reported at Google (not a startup).

New Venture Ideas

Idea #33 — Portable AI-agent trust/reputation bureau. Every agent-identity platform (NewCore, Token Security, Astrix, Offroad) is minting its own siloed agent 'trust score'; a neutral, cross-IdP reputation that travels with an agent across orgs and the MCP ecosystem is unoccupied. Directly adjacent to Helmet's agent-comms / MCP position.

Idea #34 — Server-side 'intent firewall' for agent actions. Ent evaluates intent on the endpoint, but most agent actions run server-side via APIs and MCP tool-calls where no endpoint agent sees them; an inline intent-evaluation layer at the tool-call boundary complements Ent and overlaps Helmet, Prompt Security and Pillar on runtime.

People & Thought Leaders

New entries: Zohar Alon (NewCore CEO; founder of Dome9 → Check Point); Elias Manousos (Ent CEO; founder of RiskIQ → Microsoft, Security Copilot); Dan Lorenc (Chainguard CEO; OSS supply-chain leader behind Sigstore/SLSA, now leading the Athena coalition); Shalev Hulio (Dream CEO; ex-NSO Group).

Activity: Anthropic's Project Glasswing/Mythos findings now feed the Chainguard-led Athena clearinghouse, tying Jason Clinton's team to the week's biggest open-source-defense initiative.

Helmet Watch

No Helmet Security (helmet.sh) news in-window — no new funding, hires, product launches or press identified. Helmet continues on its December 2025 \$9M seed (SYN Ventures, WhiteRabbit Ventures) and is listed as a Silver member of the Agentic AI Foundation. The week's signals run squarely through Helmet's lane: the agent-identity land-grab (NewCore), the server-side 'intent firewall' gap (Idea #34), and rising demand for MCP / agent-to-agent security.

Weekly Update — Week of 2026-06-12

Normal weekly run (7-day window, 2026-06-05 to 2026-06-12; attempt 1). A quieter week for headline mega-rounds but a defining one for narrative: Infosecurity Europe (June 2-4) crowned agentic AI security the industry's central problem, and a new stealth exit — A Security — crystallized the emerging "post-Mythos" thesis that defenders need autonomous, machine-speed exploit-path validation to keep pace with weaponized AI.

Funding Activity

2026-06-08 — A Security: emerged from stealth with \$37M from Lightspeed Venture Partners and Cyberstarts, plus angels Assaf Rappaport (Wiz CEO), Yotam Segev (Cyera CEO) and Cerca Partners. New York, founded 2025. Round type undisclosed (likely seed/seed-extension). Autonomous offensive-security and remediation platform.

2026-06-10 — Cyera: reporting reconfirmed the \$300M round at a ~\$12B valuation (total raised >\$2B, >\$150M ARR). Already logged 6/2; noted here for completeness.

2026-06-10 — Pi (pi.security): de-stealthed with a \$25M Series A (Brightmind Partners and Third Point; QP Ventures, Recursive Ventures), ~\$35M raised in total incl. a \$10M 2025 seed. Angels include CrowdStrike CEO George Kurtz and Armis co-founders Dibrov and Izrael. Founders Guy Arazi (ex-Microsoft MSRC) and Yoni Ramon (ex-Tesla). This is the public launch of the company we tracked in stealth as "Nebari" (#119) — an inferred link (shared founder + CrowdStrike accelerator + Arazi's own posts), not a rebrand Pi formally announced. Reconciled, not a new entrant. Reinforces the find-vs-fix / autonomous-remediation theme alongside A Security.

Context: Crunchbase's weekly biggest-rounds recap shows security and AI still topping the table even as overall venture pace cools; the largest recent security deals (XBOW \$120M Series C, Oasis \$120M Series B) date to March and remain out-of-window.

Exits & M&A

No new in-window AI+security exits this week.

Catch-up from Infosecurity Magazine's June 8 M&A roundup (May deals not previously tracked): Infoblox completed its acquisition of Axur (announced 5/5) — AI-powered external threat and digital-risk protection, adding DRPS and CTEM to Infoblox's Exposure Management.

SecurityScorecard completed its acquisition of Driftnet (announced 5/14) — pairing Driftnet's internet-discovery engine with SecurityScorecard's TITAN AI for threat-informed third-party risk. The roundup's through-line for 2026 holds: acquirers keep buying AI/agentive expertise in-house (Torq-Jit, Akamai-LayerX already tracked).

New Entrants

A Security (a.security) — autonomous offensive security. Founders Yossi Torati (CEO; ex-Director of Enterprise Security at Sygnia), Omer Gull and Yuval Itzchakov (ex-Check Point, Hunters, IDF Unit 8200). Offensive and defensive AI agents continuously discover, validate the exploitability of, and remediate real-world attack paths before AI-enabled attackers chain them — with scoped execution and full audit trails. Competes with XBOW, RunSybil, Horizon3.ai, Pentera and Aisy. Differentiator: closing the loop from validation to proof-of-remediation, not just flagging risk.

Theme Signals

Agentive AI security is now the industry's headline problem. At Infosecurity Europe (6/2-4), OWASP launched an Agentive Research Council and an Enterprise Adoption Maturity Model atop its Agentive Top 10; researchers demonstrated an open-weight AI worm compromising 73.8% of a simulated enterprise network.

The readiness gap is the TAM: Cisco's State of AI Security 2026 found only 29% of organizations feel prepared to secure agentive deployments while 83% plan to deploy them — a wide deploy-vs-secure gap feeding posture, guardrail and maturity-scoring tooling (Noma, Zenity, Straiker).

Prompt injection “remains unsolved” (OWASP researcher, 6/8) — a recurring reminder that the foundational agentive attack class still lacks a durable defense.

Offense and remediation are converging into a single autonomous loop. A Security's “post-Mythos” framing — validate exploitability, then remediate at machine speed — sharpens the autonomous-pentest category beyond flag-only testing and reframes exposure management as continuous proof-of-closure.

Failures & Pivots

None identified in the AI+security startup space this window.

New Venture Ideas

Idea #31 — Exploit-path validation feed for cyber-insurance underwriting. Insurers price AI-era risk on guesswork; autonomous validators (A Security, XBOW) now produce auditable proof of which exploit paths are actually open or closed. Wedge: sell that evidence to underwriters and MGAs, not CISOs. Moat: a data network effect across insured exploit-path telemetry.

Idea #32 — Agentive-AI deployment gate scored to the OWASP maturity model. With 83% deploying agents and only 29% ready, there's no standard pre-deployment readiness gate. Build the de-facto OWASP-aligned scorecard with CI/CD and IdP integrations. Competes with Noma, Zenity and Straiker on posture; differentiator is standardized maturity scoring.

People & Thought Leaders

New entry: Shlomo Kramer (Cato Networks CEO/co-founder; founder of Check Point and Imperva; prolific cyber angel) delivered the Infosecurity Europe 2026 opening keynote (6/2) on technology, investment and innovation cycles — a bellwether read on where cyber capital is heading.

Activity: Johann Rehberger (OWASP) used Infosecurity Europe (6/8) to warn that prompt injection remains unsolved, reinforcing his standing as the field's most prolific agentic-AI bug-finder.

Helmet Watch

No Helmet Security (helmet.sh) news this week — no new funding, hires, product launches or press identified in-window. Last tracked milestone remains the \$9M seed (Dec 2025; SYN Ventures, WhiteRabbit) for MCP / agent-communication security. The surrounding non-human-identity and agent-communication theme stayed hot (A Security's exploit-path angle; ongoing agent-identity land-grab).

Weekly Update — Week of 2026-06-01

A policy-heavy week. The June 2 White House executive order 'Promoting Advanced Artificial Intelligence Innovation and Security' establishes an AI cybersecurity clearinghouse and directs federal grant money toward AI vulnerability detection — landing the same day Anthropic expanded Project Glasswing (Claude Mythos) to ~150 organizations after reporting 10,000+ high-severity vulnerabilities found in its first month. On the private side, Cyera doubled its valuation to \$12B in five months, and two stealth seeds (Offroad, Willow) extended the agent-identity land-grab. No AI+security shutdowns identified in-window. No Helmet news this week.

Funding Activity

2026-06-02 — Cyera: \$300M growth round at a \$12B post-money, led by Evolution Equity Partners (Georgian, Greenoaks, Lightspeed, Sequoia, Sapphire, Redpoint, Cyberstarts, Coatue, Accel, Spark participating). Valuation has gone \$6B (Nov 25) → \$9B (Jan 26) → \$12B; TechCrunch pegs the multiple at ~80x ARR despite operating losses. Total raised now exceeds \$2B.

2026-06-04 — Offroad: \$7M seed led by Ibex Investors and Skywell Capital. New York / Tel Aviv; founded May 2025 by Dan Bendler (CEO) and Philip Shteyn (CTO, ex-Unit 8200). Agentic AI that investigates, governs, and remediates identity risk across human, machine, and AI identities.

2026-06 — Willow: \$7M seed led by Hetz Ventures, with angel backing from Wix CEO Avishai Abrahami and President Nir Zohar. Founded by ex-Wix engineers Eyal Ben Ezra, Shalev Shalit, and Idan Chetrit. Agentic access platform governing how AI agents (Claude, Cursor, ChatGPT, Codex, n8n, custom) connect to internal systems; ships a 1,000+ connector marketplace.

Exits & M&A

Catch-up (2026-05-20, missed in prior run) — Cyera acquired Genie Security for an estimated \$50M. Five-month-old, five-employee Israeli startup (Nadav Noy, Noam Dotan) doing endpoint DLP against GenAI-driven data leakage; raised only \$3M seed (Mensch Capital, Dynamic Loop; Assaf Rappaport angel). Cyera's fifth acquisition after Trail (\$162M), Ryft (~\$100M), Otterize, and Shape AI.

No new in-window exits identified this week.

New Entrants

Offroad (stealth exit 6/4) — agent-era identity resolution. Launch artifact worth noting: an audit of 2,890 public OAuth apps on Google Workspace / GitHub marketplaces found 32% carry structural exposure (over-broad scopes, AI write access, dead publisher domains); also launched ohauth.ai, a community catalog of risky OAuth apps.

Willow (stealth exit, early June) — agent access governance with a connector marketplace; cites surveys showing 79% of companies deploying agents and 65% reporting agent-related incidents in the past 12 months.

Theme Signals

US policy tailwind: the June 2 EO creates an AI cybersecurity clearinghouse (voluntary, industry + critical infrastructure), directs OMB/ONCD/CISA to find grant funds for AI vulnerability detection, and expands federal cyber hiring. Direct tailwind for AI code security, vuln remediation, and autonomous pentest categories.

Find-vs-fix gap is now the macro theme: Glasswing found 10,000+ high-severity flaws (6,202 high/critical across 1,000+ OSS projects) in a month, and Anthropic itself notes remediation remains the bottleneck. Discovery is being commoditized by frontier models; the fix side (Mobb, Pixee, Cogent, AISLE, Maze) is where scarcity now sits.

Agent identity land-grab: two same-week seeds (Offroad, Willow) plus 1Password's partnership with OpenAI to stop coding agents leaking credentials — all within days of Geordie's \$30M A.

Agent access/identity governance is the most crowded seed category in the space.

Securing AI coding agents is splitting into its own category: SecurityWeek covered the SymJack attack (turning coding agents into supply-chain delivery systems), a one-click VS Code token-theft vuln, and the launch of Edamame (catching coding agents going off the rails) in the same week.

Valuation froth watch: Cyera at ~80x ARR; capital concentrating in perceived category winners while seed-stage entry continues unabated.

Failures & Pivots

None identified in the AI+security space this window. Macro backdrop: 2026 tech layoffs reached ~142,000-150,000, but cybersecurity remains a net-hiring sector.

New Venture Ideas

Idea #29 — OAuth and integration-grant hygiene for the agent era. Continuous audit + auto-revocation/right-sizing of third-party and agent OAuth grants (Google Workspace, GitHub, Slack). Offroad's own data (32% of marketplace apps risky) proves the wedge; competes with Offroad, Push Security, and Cisco/Astrix, but none lead with marketplace-app hygiene.

Idea #30 — Remediation clearinghouse infrastructure: the fix-side of Glasswing. Auto-triage, verified patch generation, and downstream distribution for OSS maintainers and enterprises, plugged into the EO's clearinghouse. Competes with Mobb/Pixee/Cogent on the enterprise side; the OSS/clearinghouse distribution angle is unoccupied and may carry a federal anchor customer.

People & Thought Leaders

Jason Clinton / Anthropic — Project Glasswing expansion (150 orgs, 15+ countries) is the most consequential vendor-side AI security move of the week; updated notes.

New entries: Dan Bendler (Offroad CEO; published the OAuth marketplace audit + ohauth.ai) and Eyal Ben Ezra (Willow CEO, ex-Wix).

People-on-the-move: Cyera appointed Naveen Palavalli as Chief Marketing Officer.

Weekly Update — Week of 2026-05-25

A week dominated by Series A announcements in the agentic-AI security stack. Two May 28 closes — Geordie AI's \$30M (Balderton-led, largest European cyber Series A on record) and Gray Swan AI's \$40M (Wing + Madrona co-led) — bracket the agent governance and frontier-model red-team layers respectively. Two further in-window seeds, CodeIntegrity (\$5M, Syn Ventures) and Canyon Code (\$5M, Cota Capital), reinforce the emerging multi-agent runtime/orchestration category that now sits squarely between model-serving and DSPM-for-AI. No major AI+security startup shutdowns this week; one small AI-MSSP roll-up (Cycurion → Secuvant). No Helmet news in window.

Funding Activity

2026-05-28 — Geordie AI: \$30M Series A at ~\$180M post-money, led by Balderton Capital with Crosspoint Capital, General Catalyst, and Ten Eleven Ventures. London-based; positions as the largest European cybersecurity Series A on record. Total raised \$36.5M. Discovery + posture + behavior monitoring for enterprise AI agents, plus the Beam remediation suite that uses context engineering to dynamically constrain agent behavior. Geordie reports 1,300% ARR growth in the first five months of 2026 across ~30 customer environments.

2026-05-28 — Gray Swan AI: \$40M Series A co-led by Wing Venture Capital and Madrona, with Obvious Ventures, Snowflake Ventures, Hudson River Trading, Samsung Next, and Magarac Venture Partners. CMU spinout founded by Andy Zou, Matt Fredrikson, and Zico Kolter (also on the OpenAI board). Builds red-team benchmarks and adversarial tools cited in Anthropic, OpenAI, and Meta safety disclosures; moving from research-tooling to enterprise red-team services with the Series A.

2026-05-27 — CodeIntegrity: \$5M seed, Syn Ventures lead; Antler and Boost VC follow-on. Runtime control layer that translates and filters agent actions and constrains which systems and data each agent can reach. Pitched explicitly as 'permanent guardrails' for unpredictable agents.

2026-05-26 — Canyon Code: \$5M pre-seed, Cota Capital lead with Newbuild Ventures and Blackhorn Ventures. Out-of-stealth launch. Workflow intelligence layer above the model serving layer for multi-agent applications — per-app/per-user optimization, access control, policy, and update governance for enterprise agent fleets.

Exits & M&A

2026-05-22 — Cycurion (NASDAQ: CYCU) → Secuvant, LLC. Terms undisclosed; the target contributes ~\$3M annualized revenue and ~\$1.5M EBITDA for FY2026. Small AI-MSSP roll-up rather than a category-defining exit, but Cycurion is now stacking acquisitions (Halo Privacy + HavenX from May 20) to build out an AI-driven managed security portfolio. Worth tracking as a microcosm of the public-co AI-cyber roll-up pattern.

New entrants and notable launches

Three new companies enter the tracker — Gray Swan AI (#160), CodeIntegrity (#161), Canyon Code (#162). Geordie AI's existing entry (#93) is updated to Series A. Adjacent product news in the window: 7AI announced PLAID ELITE and 100 new AI jobs in Boston on May 26, highlighting >7M production investigations to date — the operating-scale signal from the largest cyber Series A in history is now visible. Sysdig's headless cloud security platform (May 6, slightly out of window) continues to read as the most pointed 'CNAPP for the agentic era' positioning so far.

Theme signals

Agentic-AI orchestration and runtime governance is hardening into its own category. Four in-window rounds (Geordie, CodeIntegrity, Canyon Code, plus Gray Swan on the offensive-test side) target distinct layers of one emerging stack: discovery → policy → runtime guardrails → workflow intelligence → adversarial validation. Existing AI-SOC and DSPM-for-AI vendors sit at the wrong layer to inspect cross-agent decision chains, which is creating real white space. European cyber capital is scaling on the AI thesis. Geordie's \$30M is the largest European cybersecurity Series A on record; Balderton leading at ~\$180M post-money signals that EU growth investors now view AI-cyber as ready for late-stage check sizes. Worth watching for a follow-on cohort (Mindgard, AdvAI, ZeroPath, Aisy) over the next two quarters. Frontier-lab proximity is becoming a moat for commercial red-teams. Gray Swan's pitch is that every major lab (Anthropic, OpenAI, Meta) already cites its work — Zico Kolter's OpenAI board seat is the most explicit version of this. Pattern repeats with Haize Labs and Mindgard; venture thesis is that distribution into Fortune 500 follows from frontier-lab trust signal.

Failures & pivots

No public AI+security startup shutdowns in window. Adjacent macro signal: Cloudflare confirmed ~1,100 AI-driven layoffs (internal AI usage reportedly up 600% in 3 months). Several mid-size cyber vendors — Kaseya, Axonius, CyberArk, At-Bay, Pentera, Arctic Wolf — also showed cuts earlier this year. The industry signal: cyber demand stays strong, but back-office cyber-vendor headcount is now being optimized by AI agents inside those same vendors.

New venture ideas

Two new ideas added to the master list this week, indexed 27 and 28.

27. Vertical workflow-intelligence layer for security agents ("Datadog for agentic SOC")

Wedge: Canyon Code positions horizontally for all multi-agent apps, but the highest-value first beachhead is the SOC. Every Tier-1 AI-SOC (Dropzone, 7AI, Prophet, Tenex.AI, Exaforce) runs multi-step agent flows that today are black boxes for the CISO and uncapped on token cost. A purpose-built observability + policy + cost-attribution layer for security-agent workflows could be 'Datadog for agentic SOC' — a neutral, cross-vendor control plane. Why now: F500 SecOps teams are running 2+ AI-SOC pilots, CFOs are asking for per-investigation cost telemetry, and ISACs want cross-tenant explainability. Risk: AI-SOC vendors bundle this themselves; differentiate by going neutral-Switzerland across vendors and integrating with FinOps tooling.

28. Agent-aware DLP for coding agents at the IDE/browser edge

Wedge: Operant Endpoint Protector, LayerX (Akamai), and Frame all moved closer to the human endpoint. There is a specific gap on developer endpoints where coding agents (Cursor, Claude Code, Cline, Windsurf) routinely upload source code, secrets, and customer data to

model providers. A secrets + PII broker purpose-built for coding agents — installed as an IDE extension and browser extension, identity-aware, MCP-server-registry aware — could pre-empt a Wiz-style enterprise wedge. Early customers: regulated industries (finance, healthcare, defense) that have outright bans on coding agents that they want to relax safely. Risk: Anthropic / OpenAI / Cursor ship native enterprise DLP; differentiate by being multi-agent, deeply MCP-aware, and air-gap deployable.

People & thought leaders

Henry Comfort (ex-Darktrace) and Benji Weber (ex-Snyk), Geordie AI co-founders — added to the People & Thought Leaders tab on the Series A press cycle. Comfort is the public-facing CEO; Weber is engineering lead on Beam, the agent remediation module using context engineering to dynamically constrain agent behavior.

Andy Zou and Zico Kolter, Gray Swan AI co-founders — added to the People tab. Zou is the original author of the GCG jailbreak paper and continues to publish frontier red-team benchmarks. Kolter is CMU faculty and an OpenAI board member overseeing safety. Their Series A makes Gray Swan the clearest example of an academic-credibility commercial red-team play.

Simon Willison published 'Hardening Firefox with Claude Mythos Preview' (May 17) — a case study of Mozilla using preview model access to find and fix hundreds of Firefox vulnerabilities. Reinforces his thesis that frontier coding models have become credible junior vulnerability researchers.

Ross Haleliuk published 'Cybersecurity is really boring' (May 19) — argues the most effective security is unglamorous (patching, logs, least privilege, segmentation) and explicitly ties to agentic AI: 'AI agents amplify every bad permission model at machine speed.' Widely reshared this week.

Helmet watch

No new Helmet press cycle in window. Helmet remains tracked at \$9M seed (SYN Ventures, WhiteRabbit Ventures), MCP / agentic-AI communications security, AAIF Silver Member.

Notable adjacent signal: Syn Ventures (a Helmet investor) also led CodeIntegrity's seed this week — Syn is increasingly the institutional anchor for the agentic-AI runtime category.

Index hygiene

Companies tab now 162 entries (up from 159). Funding Rounds: 74 entries (up from 70). Exits & M&A: 35 entries (up from 34). Themes: 20 entries (up from 19; added 'agentic-AI orchestration & runtime governance'). Startup Ideas: 28 entries (up from 26). People & Thought Leaders: 35 entries (up from 31). Sources: 108 entries (up from 95). One Weekly Log row added (2026-05-25).

Weekly Update — Week of 2026-05-18

A relatively quiet week against the May 18 deep-dive backdrop, but two category-defining moves: Akamai bought LayerX for ~\$205M to plant a flag in browser-layer AI usage control, and Socket crossed the unicorn threshold on a \$60M Series C at \$1B, validating AI-grade software-supply-chain security as platform-tier. Ocean Security joined the AI-native email security cohort with a \$20M Series A from Lightspeed. The run also picks up two items that slipped past the May 18 deep dive: Exaforce's \$125M Series B at \$725M (May 12) and OpenAI's March acquisition of Promptfoo.

Funding Activity

2026-05-20 — Socket: \$60M Series C at \$1B valuation, led by Thrive Capital with Andreessen Horowitz, Abstract Ventures, and Capital One Ventures. Cumulative funding \$125M. Open-source dependency analysis platform now expanding into Socket Firewall, certified patches, browser extensions, code-editor extensions, AI tools and MCP servers. Customers include Anthropic, xAI, Replit, Cursor, Figma, Vercel, Gusto, Mercado Libre, and Cribl plus F100 financial-services and media firms. First pure-play AI code/supply-chain security vendor to cross unicorn on the open-source-dependency thesis (Chainguard/Endor sit adjacent on container and SCA respectively).

2026-05-19 — Ocean Security: \$20M Series A led by Lightspeed Venture Partners with Picture Capital and Cerca Partners. Stacked with an \$8M 2024 seed = \$28M cumulative. Founders Shay Schwartz (CEO) and Oran Moyal (CTO) are both IDF Intelligence Corps veterans (Schwartz also contributed to Iron Dome research). The product runs AI agents that inspect every message in real time, evaluating sender identity, content, context, and links to detect AI-driven phishing — replacing legacy pattern-matching defences. Already 7-figure ARR; customers include Kayak and Kingston; over 1B emails processed. Strong angel stack: Wiz CEO Assaf Rappaport, Armis founders Yevgeny Dibrov & Nadir Izrael, Island founders Mike Fey & Dan Amiga, Transmit founders Rakesh Loonkar & Mickey Boodaei. Adds to the rapidly-densifying AI-native email cohort alongside Abnormal, StrongestLayer, AegisAI, and Material.

2026-05-12 (catch-up — missed in May 18 deep dive) — Exaforce: \$125M Series B at \$725M post-money valuation. Investors include HarbourVest, Peak XV, Mayfield, Khosla Ventures, and Seligman Ventures. Came roughly a year after the \$75M Series A; cumulative funding ~\$200M. Autonomous “Exabot” agents triage and stop threats in real time. Sits in the Tenex.AI / Dropzone / Prophet / Conifers agentic-SOC bracket and now ranks among the most-funded pure-plays in U.S. AI-SOC.

2026-03-18 (catch-up — missed in May 18 deep dive) — Manifold Security: \$8M Seed led by Costanoa Ventures with Cherry, Rain Capital, and Modern Technical Fund; notable angels include Joe Sullivan (ex-Uber CSO) and Vijay Bolina (ex-Google DeepMind CISO). AI Detection & Response (AIDR) for autonomous AI agents on enterprise endpoints. Manifold also pushed a Manifest catalog expansion (May 12, 2026) scoring 7,700 MCP servers — positioning the product as both a runtime defender and a supply-chain inventory plane for agent-security teams. Adjacent to Helmet (MCP server identity/auth) and Capsule (agent-side controls).

Exits & M&A

2026-05-15 — Akamai → LayerX, ~\$205M cash, expected to close Q3 2026. LayerX is a Tel Aviv-founded secure-enterprise-browser extension + AI usage governance vendor (will not force a managed browser). The acquisition extends Akamai’s Zero Trust portfolio into the browser layer where employees actually interact with consumer GenAI and agentic AI tools. Akamai expects ~\$0.12 non-GAAP EPS dilution in FY26; LayerX ARR projected ~\$10M by year-end. Strategic signal: browser-layer AI usage control is now its own acquisition category — this is the third browser-AI security exit in two quarters (Zscaler → SquareX, CrowdStrike → Seraphic preceded). Microsoft Edge for Business / Google for Workspace will likely face pressure to respond inside Q3.

2026-03-09 (catch-up — missed in prior runs) — OpenAI → Promptfoo, undisclosed terms. Promptfoo’s last priced round was a \$18.4M Series A in July 2025 (Insight Partners-led, \$86M post-money); cumulative \$23M. The open-source LLM red-team / eval library is used by 125K+ developers and 25%+ of the Fortune 500. Tech to be folded into the OpenAI Frontier enterprise agent platform. This is OpenAI’s clearest move yet into AI application security and parallels

Anthropic's investments in defender-side capability (Mythos, Claude Security GA). Promptfoo founders Ian Webster and Michael D'Angelo expected to join OpenAI's security org.

New entrants and notable launches

Three companies added to the tracker this week: LayerX (#157), Ocean Security (#158), and Manifold Security (#159). The May 12 Manifold "Manifest" catalog expansion to 7,700 MCP servers is the largest single shipped artifact for agent-security teams since the Astrix announcement — it gives blue-team practitioners a public, searchable map of the MCP supply chain. Socket's expansion into MCP server analysis pushes the company directly into the agentic supply-chain space too, blurring the line between "code supply chain" and "agent supply chain."

Theme signals

Browser-layer AI usage control hardens as its own category. Akamai → LayerX is the third browser-AI security exit in two quarters (after Zscaler → SquareX in Feb and CrowdStrike → Seraphic in Jan). The browser is now an accepted enforcement chokepoint for AI DLP, paste-blocking, content classification, and agent oversight — sitting in a sweet spot between ZTNA/SASE control planes and CASB/DSPM data planes. The new theme has been added to the Themes tab. Watch Microsoft (Edge for Business + Defender for Cloud Apps) and Google (Chrome Enterprise + Gemini for Workspace) for their counters; Island and Talon (PANW) are the obvious next targets if a platform vendor wants the managed-browser version. AI-MDR is the dominant commercial wrapper for AI-SOC tech. Between Tenex.AI (\$1B val), Exaforce (\$725M), Dropzone, Prophet, Andesite, and 7AI, AI-SOC vendors are increasingly selling humans-plus-agents as a service rather than seat-licensed software. F500 buyers want "someone to call when the agent fails" — particularly in regulated industries. This is a structural shift from the SaaS-only thesis investors had in 2024; expect renewal economics to reflect MDR-style ARR.

Software supply chain crosses unicorn threshold on the AI thesis. Socket's \$1B Series C, with explicit MCP-server and AI-tools expansion plans, places SCA / dependency analysis squarely inside the AI-security perimeter. Chainguard, Endor Labs, Apiiro, and Semgrep should expect comparable valuation marks if buyers continue to grade vendors on agentic-supply-chain readiness.

AI-native email security is densifying fast. Ocean Security joins Abnormal, Material, StrongestLayer, AegisAI and Frame at this layer. Differentiation is narrowing to (a) agent architecture quality, (b) integration with collaboration suites (Slack, Teams), and (c) post-detection workflow (auto-revoke, identity verification, downstream NHI rotation). Expect at least one more \$50M+ round in the cohort by end of Q3.

Failures & pivots

No public AI+security startup shutdowns in the window. Adjacent: Cloudflare confirmed an additional ~1,100 AI-driven layoffs this week and Meta's ongoing 8,000-person workforce reduction touches integrity and security teams. These continue to be platform-vendor cost-base resets rather than AI-native-cyber distress, but they keep the "AI replaces the SOC analyst, then replaces the AI SOC builder" narrative alive at the CFO level — relevant for AI-MDR pricing power and for incumbent vendor pricing pressure.

New venture ideas

Two new ideas added to the master list this week, indexed 25 and 26.

25. Pure-play AI-MDR for the agentic stack

Wedge: Tenex.AI and Exaforce prove the AI-MDR commercial model at unicorn / near-unicorn marks, but neither is positioned specifically for the agentic stack itself — protecting MCP servers, agent runtimes, tool-use governance — as a managed service. SOCs increasingly run agents (Microsoft MDASH, AWS Rex, Cisco DefenseClaw) and need 24/7 humans-in-the-loop monitoring those agents. Hybrid: own-IP agentic detection content + thin services layer + Helmet/Capsule/Manifold-class tooling resold. Buyer is the F500 CISO already in late-stage Astrix/Microsoft Agent 365 / Anthropic Skills POCs. Risk: incumbent MDRs (Arctic Wolf, Red Canary, Expel) add agent monitoring; differentiation has to be agentic-stack-native, not a general SOC bolt-on.

26. AI usage observability for the audit committee, regulators, and cyber insurers

Wedge: Akamai paid \$205M for LayerX at ~\$10M ARR — the AI-usage telemetry is the asset, not the browser. But the data is consumed by CISOs; the audit committee, the SEC-cyber-rules drafter, the EU AI Act enforcement officer, and the cyber-insurance underwriter all want a different abstraction: “show me where AI created material risk this quarter and prove who approved it.” Today this lives in Word docs and Excel. A vertical product would aggregate from existing usage controls (LayerX/Akamai, Harmonic, Nightfall, Microsoft Defender for Cloud Apps, Zscaler/Netskope), generate audit-grade attestations, and map to EU AI Act + NIST AI RMF + NYDFS + Coalition/Resilience underwriting questionnaires. Distribution via Big-4 audit firms. Adjacent to Lema AI, Compliance, Vanta/Drata.

People & thought leaders

Shay Schwartz (CEO) and Oran Moyal (CTO), Ocean Security — both IDF Intelligence Corps veterans; Schwartz also contributed to Iron Dome research. Added to the People tab. Ian Webster, Promptfoo co-founder/CEO — moves to OpenAI's security org with the March 2026 acquisition. One of the most-cited operators on practical LLM red-teaming / eval. Worth tracking inside OpenAI's defender-side capability roadmap.

Helmet watch

No new Helmet press cycle in the window. Helmet remains tracked at \$9M seed (SYN Ventures, WhiteRabbit Ventures), MCP / agentic-AI communications security, team building toward 10. Manifold's Manifest expansion and Socket's MCP-server scope are the two most directly competitive moves of the week — both broaden the “MCP supply-chain” coverage that Helmet sells against; differentiation pressure increases on the identity / policy enforcement axis Helmet leads on.

Index hygiene

Companies tab now has 159 entries (up from 156). Funding Rounds: 70 entries (up from 66). Exits & M&A: 34 entries (up from 32). Themes: 19 entries (up from 18). Startup Ideas: 26 entries (up from 24). People & Thought Leaders: 31 entries (up from 28). Sources: 95 entries (up from 82). Weekly Log: row added for week of 2026-05-18.

Weekly Update — Week of 2026-05-11

This week extended two of the strongest patterns of the year: agentic SOC funding stays at scale and frontier-lab + platform vendor defender-AI capability moves from research into shipped product. The week also marks the first \$50M check into a pure "human risk security" startup and the largest set of AI-driven layoffs at security-adjacent vendors to date.

Funding Activity

2026-05-12 — Exaforce: \$125M Series B at \$725M valuation, led by HarbourVest with Peak XV, Mayfield, Khosla Ventures, Seligman Ventures, and AICONIC. Three-year-old AI-SOC company combines a real-time security knowledge graph with autonomous "Exabot" AI agents and an MDR service; total funding now \$200M. Customers include Replit and Guardant Health. Capital earmarked for multi-model AI development and GTM expansion in Japan and Europe. Sits in the same agentic-SOC bracket as Prophet, Dropzone, 7AI (\$130M Series A late 2025), and Conifers.

2026-05-11 — Frame Security: \$50M out of stealth, led by Index Ventures with Team8 and Picture Capital. Founded by Tal Shlomo (ex-Wiz early employee) and Sharon Shmueli (ex-Israeli Unit 8200). AI-driven platform defending against AI-generated social engineering and deepfake attacks via auto-generated role-personalized simulations and adaptive training. Already deployed at Louis Dreyfus, AlphaSense, and Rockefeller Capital Management. Defines a new "human risk security" category — directly competitive with KnowBe4-class incumbents but rebuilt on AI-vs-AI simulation generation.

Exits & M&A

No new AI+security M&A closed or announced in the 5/8–5/15 window. The pipeline from prior weeks (Cisco → Astrix intent 5/4; ServiceNow → Armis closed 4/20; Cyera → Ryft closed 4/23) remains the active set. The 247WallSt list out 2026-05-12 names SentinelOne and four other pure-plays as the most-discussed acquisition candidates for the remainder of 2026.

New entrants and notable launches

Microsoft announced MDASH (Microsoft Defender Agentic Scanning Harness) on 2026-05-12 — a multi-model agentic system that orchestrates more than 100 specialized AI agents across frontier and distilled models to discover, validate, and prove exploitable vulnerabilities end-to-end. MDASH found 16 new vulnerabilities across the Windows TCP/IP stack, IKEEXT IPsec service, HTTP.sys, Netlogon, DNS resolution, and the Telnet client — including four Critical remote-code-execution flaws (notably CVE-2026-33827 in the IPv4 stack and CVE-2026-33824 in IKEEXT, the latter affecting RRAS VPN, DirectAccess, and Always-On VPN deployments). It scored 88.45% on the public CyberGym benchmark, 96% recall on five years of historical MSRC cases for clfs.sys, and 100% recall in tcpip.sys. Microsoft is now on pace to break its annual CVE record. MDASH is in limited private preview with select customers. This is a peer to OpenAI GPT-5.5-Cyber (5/7) and Anthropic Mythos Preview (4/30); the defender-side capability uplift is now firmly inside platform vendors.

Adjacent: Anthropic upgraded Claude Security beta to GA on Claude Opus 4.7 with multi-stage validation pipeline, confidence ratings on every finding, and CrowdStrike / Microsoft Security / Palo Alto Networks / SentinelOne / TrendAI / Wiz all integrating Opus 4.7 into their existing platforms. Snyk announced Claude embedded in its AI Security Fabric. Frontier-lab cyber tiers are no longer experimental — they are now the underlying engine for the entire major-vendor stack.

Companies tab additions this week: Exaforce (#117), Frame Security (#118).

Theme signals

Human risk security becomes a budget line. Frame Security's \$50M from Index/Team8/Picture is the first nine-figure-anchored check into a pure-play AI human-risk vendor. The wedge: legacy security-awareness training is one-size-fits-all annual content; the new category generates role-personalized AI deepfake/vishing/BEC simulations on demand and measures behavior change. Adjacent to Adaptive Security (\$43M; a16z + OpenAI Startup Fund) and to AI-native email (Abnormal, Material). New theme added to tracker.

Defender AI vulnerability discovery becomes table-stakes. With Microsoft MDASH joining Anthropic Mythos, OpenAI GPT-5.5-Cyber, and XBOW, every platform vendor now needs either a native capability or a binding to a frontier-lab cyber tier. The differentiating axis shifts from "can your AI find bugs?" to "how many, how fast, with what false-positive rate, integrated with what workflow?". Microsoft's 88.45% CyberGym score is the new public bar. New theme added. Open-source agent orchestration narrows proprietary moats. AWS Rex (5/4, open-sourced container for hallucinated-code / prompt-injection / over-eager-agent containment), Cisco DefenseClaw (RSAC 2026, open-source secure-agent framework), and Cisco's May 1 open-sourcing of the Model Provenance Kit collectively suggest the security-agent control plane will not be a single proprietary winner. This is a positive signal for the venture idea pipeline: orchestration + skills + provenance can be vendor-neutral.

Agentic-SOC capital concentration continues. Exaforce (\$125M Series B at \$725M) joins 7AI (\$130M Series A late 2025), Prophet, Dropzone (\$37M Series B Jul 2025), Conifers, Andesite, Simbian, and Radiant as well-funded agentic-SOC plays. The category has crossed the threshold where Sequoia/Khosla/Index/Insight are repeat backers. Buyer signal: F500 SOC pilots are converting; the cost-per-alert math is showing up in renewal cycles.

Cisco-Astrix officially became a category-defining moment. Multiple analyses this week (Futurum, Network World, Security Boulevard) frame the \$400M Astrix deal as the proof point that non-human / agentic identity is now strategic for every platform vendor — meaning Microsoft, Google Cloud, and Palo Alto will need a comparable buy or build in the next two quarters.

Failures & pivots

No notable pure-play AI+security startup shutdowns or layoffs in the window. Adjacent: Cloudflare announced 1,100 layoffs on 2026-05-08 — the first mass layoff in its 16-year history — and explicitly attributed the cuts to AI making roles obsolete (AI usage internally up 600% in three months). Cisco announced ~4,000 layoffs (~5% of workforce) on 2026-05-14 alongside record quarterly revenue, redirecting capital to AI and cybersecurity investment. Both are security-adjacent vendor moves and signal an AI-driven cost-base reset across the major-platform stack; they do not yet imply distress at AI-native cyber startups, but they do reduce hiring competition for AI/security engineers and may accelerate consolidation at the platform tier.

New venture ideas

Two new ideas added to the master list this week, indexed 23 and 24 in the tracker.

23. Mid-market human-risk security platform delivered via MSPs

Wedge: Frame validates the buyer at the enterprise tier; mid-market healthcare provider groups, RIAs / wealth managers, regional manufacturers, hospitality, and professional-services firms face the same AI-deepfake CFO calls, vishing, and BEC attacks but cannot afford an enterprise

human-risk stack. A productized, MSP-delivered, role-personalized platform with vertical-tuned simulations and insurer-aligned reporting captures the segment. Frame will not serve directly. Differentiation: MSP-native economics (per-seat, RMM-integrated), vertical content libraries (HHS-auditor calls, custodian wire instructions, ERP cutover requests), and binder-ready evidence for cyber-insurance renewal.

24. Vendor-neutral defender-AI agent orchestration platform ("DefenseGraph")

Wedge: MDASH proves the architecture — orchestrating 100+ specialized agents beats single-shot LLM queries — but locks the IP inside Microsoft. AWS Rex and Cisco DefenseClaw signal the industry will not converge on a single proprietary winner. A vendor-neutral framework + skills marketplace for defender agents (vuln triage, IR investigation, threat-intel correlation, detection authoring) spanning Anthropic, OpenAI, Google, and open models. Pre-mapped to NIST AI RMF + EU AI Act + DoD IL5/6; partnerships with verified-defender tiers across labs. Adjacent to Ideas #20 (model provenance SaaS) and #21 (skill inspection platform) — same team could build the full agentic-security control plane.

Index hygiene

Companies tab now has 118 entries (up from 116). Funding Rounds: 40 entries (up from 38). Exits & M&A: unchanged at 25. Themes: 16 entries (up from 14). Startup Ideas: 24 entries (up from 22). Sources: 61 entries (up from 51). Weekly Log: row added for week of 2026-05-11. Weekly Update — Mid-week supplement 2026-05-08 (catch-up additions)
Coverage window: standing 7-day window plus targeted catch-up of items missed by prior runs (Mar 10 — May 8). This is a mid-week supplemental run of the standing weekly task; the same-day run-2 entry above remains the canonical 5/4–5/8 weekly update. Items here are dominantly retroactive corrections to the tracker, plus a new theme triggered by 5/7 incumbent moves.

Funding activity (catch-up)

2026-03-10 — Kai emerged from stealth with \$125M (combined seed + Series A) led by Evolution Equity Partners with N47 and strategic co-investors. Founders Galina Antova (CEO, ex-Claroty co-founder) and Damiano Bolzoni (CTO, ex-SecurityMatters, sold to Forescout 2018) are building an agentic AI cybersecurity platform that bridges IT and OT — a Claroty-class buyer with an agentic SOC layer on top. Reported 7-figure bookings within ten months across energy, pharma, automotive, and hospitality. This was missed by prior weekly runs and is now added to Companies and Funding Rounds.

2026-03-19 — Oasis Security closed a \$120M Series B led by Craft Ventures with Cyberstarts, Sequoia, and Accel; total funding now \$195M. Oasis introduced an enterprise-grade Agentic Access Management (AAM) platform with intent-based, just-in-time access for AI agents and machines, positioning the company as the pure-play counter-narrative to the Cisco-Astrix consolidation. The Companies row has been updated from Series A / \$40M to Series B / \$195M and the round added to Funding Rounds.

Exits & M&A (catch-up)

2026-04-13 — Fortreum (Gryphon Investors-backed) acquired Kovr.AI, a FedRAMP Moderate-authorized AI-native compliance platform with Agent Artemis (agentic AI in a Zero-Data-Retention environment) and patented "build once, map anywhere" control mapping across FedRAMP, CMMC 2.0, DOD SRG, NIST CSF 2.0, and GovRAMP. Kovr is already deployed at the U.S. Air Force, Space Force, and Accenture Federal Services. Combined with Fortreum's C3PAO and FedRAMP assessor practice, this creates an end-to-end "agentic GRC + independent assessment" stack — an early example of services consolidators productizing agentic-AI compliance. Terms not disclosed.

2026-04-23 — Cyera acquired Israeli startup Ryft Data for an estimated \$100-130M. Ryft's automated data-lake platform gives enterprises traceable data access for AI agents — Cyera's

fourth acquisition in five years and a deliberate add to its \$9B-valuation DSPM platform aimed squarely at the data-layer of agentic deployments.

Incumbent product launches and frontier-lab moves (5/7)

OpenAI launched GPT-5.5-Cyber on 2026-05-07 to vetted defenders under its Trusted Access for Cyber program — a fewer-guardrails variant of GPT-5.5 tuned for vulnerability identification and triage, patch validation, malware analysis, simulation, and proof-of-concept exploit writing (credential theft and net-new malware authoring remain blocked). This is one month after Anthropic's Claude Mythos Preview and Project Glasswing consortium, and is the second frontier-lab "trusted-tier" defender release.

Cognizant launched Cognizant Secure AI Services on 2026-05-07 — a systems-integrator productization of "secure, govern, and scale agentic AI" across enterprise operations. SI/MSSP entry into the agentic-AI-security market is now mainstream, putting a price-floor on what pure-play startups can charge while expanding the buyer education curve.

Shutdowns / distress

No notable AI+security pure-play shutdowns or layoffs surfaced in this window beyond what was logged earlier in the week. Broader industry-wide tech layoff numbers continue to climb but the AI+security category remains protected at incumbents.

Themes — what shifted this week

New theme added to tracker: Frontier-lab "Trusted Access for Cyber" tier. OpenAI's GPT-5.5-Cyber on 5/7 plus Anthropic's Mythos Preview / Project Glasswing on 4/30 establish a new vendor pattern — frontier labs ship special, less-restrictive variants of flagship models to vetted defender teams via verified-access programs. This is a meaningful capability-uplift channel matching offensive-AI risk and creates a new product surface for SOC, IR, and threat-intel vendors that integrate these tiers (CrowdStrike, PANW, SentinelOne, Trend Micro, Wiz are early integrators). Open questions: entitlement-management complexity for SOCs spanning multi-lab trust tiers; insurance / regulatory mandates for vetted-defender access; long-tail liability under EU AI Act and DoD frameworks.

Reinforcement of prior themes: NHI / agentic IAM is now a platform-vendor must-have (Astrix ↔ Oasis dyad — strategic acquirer + pure-play scale-up); IT-OT-converged agentic security re-validated by Kai's capital and credentials; AI-native FedRAMP / agentic GRC validated commercially via Fortreum-Kovr; data-layer-of-agents validated by Cyera-Ryft.

New idea added

We are adding one new idea this week, indexed 22 in the tracker.

22. Mid-market IT/OT-converged agentic SOC

Wedge: Kai's \$125M is aimed at F500 IT/OT convergence (Claroty-class buyer). Mid-market manufacturers, regional utilities, and hospital systems run the same OT footprint without the budget for Claroty + Forescout + a SOC vendor. A productized agentic platform that bridges IT and OT for the mid-market — preset playbooks, MSP-delivered, insurer-aligned — can claim the segment Kai will not. Why now: OT-targeted ransomware is up sharply in 2026; cyber-insurance carriers are underwriting OT controls; the Five Eyes May 2026 joint guidance flags agentic AI risk in critical infrastructure. Moat / differentiation: vertical-specific OT detection content; channel-led GTM via MSPs and cyber-insurance carriers; pre-mapped CMMC / NIS2 / IEC 62443 control overlays; bundled IR retainer.

Index hygiene

Companies: 115 → 116 (Kai added; Oasis row updated to Series B / \$195M total). Funding Rounds: 37 → 39 (Kai \$125M, Oasis \$120M Series B). Exits & M&A: 23 → 25 (Fortreum-Kovr.AI, Cyera-Ryft). Themes: 14 → 15 (Trusted Access for Cyber tier). Startup Ideas: 21 → 22. Sources: 42 → 52. Weekly Log: a third 2026-05-08 row added marked '(run 3)' for the catch-up.

Weekly Update — Week of 2026-05-08

Coverage window: 2026-05-04 to 2026-05-08 (mid-week supplement; previous successful run was 2026-05-04 12:20 UTC). The headline this week is Cisco's announcement of intent to acquire Astrix Security for ~\$400M on 2026-05-04 — the largest pure-play non-human identity (NHI) / AI-agent security exit to date and a defining marker that platform vendors now treat agentic IAM as strategic infrastructure rather than a feature. Microsoft Agent 365 went GA on 2026-05-01 as a productized agent control plane, and incumbents Netskope and ThreatBook shipped agentic SOC offerings within days of each other. The week was light on net-new pure AI+security funding rounds; capital that did move went to space security (True Anomaly \$600M Series D) and consumer privacy (Cape \$100M Series C at \$900M) — adjacent but outside our pure-play scope.

Funding activity

No new pure AI+security venture rounds of material size surfaced in the 2026-05-04 to 2026-05-08 window. Crunchbase's top-10 weekly U.S. roundup for the period was led by Anthropic's reported \$5B Amazon investment (general AI infra, not security), True Anomaly \$600M Series D (space defense), Reliable Robotics \$160M (autonomy), Cape \$100M Series C (privacy-focused mobile network — privacy-adjacent), and AcuityMD \$80M (medtech AI). Sierra raised an additional \$950M at a \$15B valuation on 2026-05-04 (general agent infra). Quantitatively, cybersecurity sector equity growth funding remains up materially YoY (Capstone Partners May 2026 update: +189.7% YoY through YTD 2026). Qualitatively, the pure AI+security category was on pause this week — capital appears to have rotated toward defense/space and adjacent privacy plays while the agentic IAM category digests the Cisco-Astrix benchmark.

Exits & M&A

- 2026-05-04 — Cisco announces intent to acquire Astrix Security for ~\$400M (per Calcalist, Globes, Times of Israel). Astrix (founded 2021 by Unit 8200 vets Alon Jackson + Idan Gour) had raised >\$85M from Bessemer / CRV / Menlo Ventures and was last marked at ~\$200M. The deal will fold Astrix into Cisco Identity Intelligence, Secure Access, and Duo to discover, govern, and secure AI agents and non-human identities. Strategic significance: this is the largest pure-play NHI exit to date, and the first \$400M+ outcome explicitly priced on the "secure the agentic workforce" thesis. Implication: platform vendors (Cisco, Palo Alto, CrowdStrike, Microsoft, ServiceNow, Okta) will likely race to lock down their NHI / agent-IAM stack — expect Oasis, Token, Astrix-class peers to see follow-on platform-vendor interest within 1–2 quarters.
- 2026-05-07 — NextPoint Group (Godspeed Capital portfolio) acquires UScontracting to expand intelligence-focused capabilities (cybersecurity, signals intelligence, data analysis) for U.S. national-security customers. Terms not disclosed. Significance: govcon roll-up rather than a pure AI+security signal, but reinforces the broader IC / defense pull on cyber consolidation.

New entrants and incumbent product launches

- Microsoft Agent 365 — GA on 2026-05-01. Microsoft's control plane for AI agents went general availability for Commercial customers at \$15/user/month (or bundled in the new Microsoft 365 E7 "Frontier Suite"). Agent 365 assigns each AI agent its own Entra identity, applies Purview labels, governs runtime behavior through Defender, extends

Intune device management to local agents on Windows, and ships registry sync with AWS Bedrock and Google Cloud Gemini Enterprise Agent Platform (multicloud governance). Implication: every AI-agent-IAM startup now competes directly with a productized Microsoft offering at \$15/user; differentiation must come from non-Microsoft estate coverage, deeper runtime, or skill / tool inspection that Microsoft does not yet ship.

- Netskope AgentSkoPe — launched 2026-05-06. New AI agent foundation inside Netskope One Platform that automates SOC / NOC alert triage, investigations, policy support, configuration checks, troubleshooting, and risk queries. Significance: another incumbent (Netskope) joins the agentic SOC race; pressure on pure-play AI-SOC startups (Dropzone, Prophet, Radiant, Simbian, 7AI) intensifies — incumbents bundle agents into existing platforms, removing buying friction.
- ThreatBook Flocks + SafeSkill — launched 2026-05-07. Flocks is an AI-native SecOps platform; SafeSkill inspects and evaluates AI-agent skills (pre-import inspection, marketplace filtering, download scanning, inventory remediation to detect tampering or hidden risks). Significance: SafeSkill is the first material productization of a new category — AI agent skill / tool marketplace security — that we are now adding to our themes list. ThreatBook is a Beijing-based incumbent; the equivalent net-new Western startup play is wide open.

Shutdowns / distress

No notable AI+security pure-play shutdowns or layoffs surfaced this week. Broader industry layoffs continue (Q1 2026 tech industry cuts ~78K per Tom's Hardware / CNBC, with ~48% AI-driven), but security organizations remain protected at incumbents — Oracle's 30K cuts explicitly carved out security, AI data center, and database engineering. Anthropic's exclusion from the Pentagon IL6/IL7 list (logged in last week's update) remains the most material commercial signal in adjacent territory.

Themes — what shifted this week

- Non-human identity / agentic IAM is now strategic must-have. Cisco's \$400M Astrix bid is the explicit benchmark. Expect Palo Alto Networks, CrowdStrike, Microsoft (likely organic), ServiceNow, and Okta to shop the remaining Astrix-class portfolio (Oasis, Token, Clutch, Entro, Aembit, Britive, Andromeda) within two quarters. NHI / agent-IAM follows the same path DSPM took in 2024: the moment one platform vendor pays a strategic premium, the rest move to defend.
- Agent control planes go GA. Microsoft Agent 365 GA on 5/1 productizes the "secure the agentic workforce" pitch into a \$15/user-per-month SKU at scale. Mid-market and SMB segments will adopt the Microsoft default; the pure-play opportunity narrows to (a) non-Microsoft-estate coverage, (b) advanced runtime / skill inspection, and (c) regulated-industry differentiation.
- AI agent skill / tool integrity emerges as a distinct sub-category. ThreatBook SafeSkill (5/7) is the first incumbent productization. MCP servers, agent skills, and agent tool marketplaces have grown ~10x QoQ; multiple poisoned-skill / prompt-injection-via-tool incidents reported in Q1 2026. We are adding this as theme #13 in the tracker.
- Agentic SOC race firmly favors incumbents. Netskope, Arctic Wolf, Microsoft Sentinel, ThreatBook, and Cisco have all launched or expanded agent-fleet offerings within the past 60 days. Pure-play AI-SOC startups (Dropzone, Prophet, Radiant, Simbian, 7AI, Crogl, Andesite) need either unique data (proprietary telemetry / partnerships) or unique

distribution (MSSP-first, regulated-industry-first) to defend their wedge. Expect the first AI-SOC pure-play layoff or strategic-alternatives announcement within two quarters.

- Pure AI+security funding gravity took a breath this week. After ~\$3.6B in cumulative agentic-AI-security funding in the RSAC / Q1 2026 cluster, the May 4–8 window saw capital rotate to space security and AI privacy. This is likely a near-term pause rather than a topping pattern; expect renewed AI+security pure-play activity around Black Hat USA (August) and the first independent AI-firewall validation results.

New ideas added to the pursue list

We are adding one new idea this week, indexed 21 in the tracker.

21. Vendor-neutral AI agent skill & tool inspection platform

Wedge: as agents proliferate via MCP servers, agent skills, and third-party tools, every enterprise needs to inspect, fingerprint, and gate these capability artifacts before import — and continuously verify them at runtime. ThreatBook SafeSkill (5/7) validates the buyer pain in the regional / incumbent slice; the bigger play is a vendor-neutral platform that spans MCP, OpenAI Skills, Anthropic Skills, Google Vertex tools, LangGraph, CrewAI, and homegrown agent frameworks. Why now: SafeSkill ships, Cisco-Astrix prices the adjacent identity slice at \$400M, and MCP server count is up 10x QoQ. Early customers: F500 platform-engineering teams, AI / ML governance leads in regulated industries, MSSPs. Moat: multi-vendor coverage breadth, library of skill fingerprints + adversarial test harness, integrations with agent platforms, enterprise GRC mappings. Adjacency to Astrix (NHI for agents) is the natural bundle: identity + skill-integrity is the agentic security control plane.

Index hygiene

Companies tab unchanged at 115 entries (Astrix updated to "Acquired" status). Funding Rounds: 36 (no qualifying new rounds in window). Exits & M&A: 22 (up from 21 — Cisco/Astrix added). Themes: 13 (up from 12 — added "AI agent skill / tool marketplace security"). Startup Ideas: 21 (up from 20). Sources: 41 (up from 31 — 10 new sources for this run).

Weekly Update — Week of 2026-05-04

Coverage window: 2026-04-28 to 2026-05-04 (normal weekly cadence; previous successful run was 2026-04-27). The headline this week is that "AI Gateway" has been formalized as a category by Palo Alto Networks, frontier AI labs are now central to U.S. cyber policy, and defenders shipped multiple agentic-AI countermeasures the same week Anthropic disclosed Mythos's near-instant exploitation capability. Pure AI+security venture financing was quiet in the U.S. window; the activity moved to platform M&A, partnerships, and product launches.

Headline events

On April 30, Palo Alto Networks announced its intent to acquire Portkey, an AI Gateway provider already processing trillions of agent tokens per month, and used the deal to declare the

AI Gateway the “mission-critical control plane” for autonomous agents. Portkey will become the AI Gateway for Prisma AIRS, layering identity, semantic routing, audit, and policy controls over every agent-to-agent and agent-to-tool interaction. Financial terms were not disclosed; the transaction is expected to close in PANW’s FQ4 FY26. The same day, PANW announced a strategic partnership with Armadin, the new offensive-security company founded by ex-Mandiant CEO Kevin Mandia, to embed Armadin’s autonomous AI attack swarm into Unit 42’s Frontier AI Defense service for continuous external attack validation. Combined with the previously closed Koi acquisition (April 14, ~\$400M) and the Chronosphere close (January 29, \$3.35B), PANW is now the most aggressive platform consolidator in agentic AI security, spanning gateway, endpoint, observability, and offensive validation.

The other major event from the prior week that needed to be captured for completeness: ServiceNow closed its \$7.75B acquisition of Armis on April 20, the largest pure cybersecurity exit since the Google–Wiz close, and the largest AI-cyber deal centered on OT/IoT exposure management. ServiceNow leadership publicly framed the combination as a path to “the largest cyber business in the world.”

On the policy and frontier-model front: Axios reported that OpenAI and Anthropic each held closed-door classified briefings with the House Homeland Security Committee on the cyber-offensive capabilities of their newest models. Anthropic delayed Mythos Preview because the model can rapidly identify and exploit vulnerabilities; OpenAI moved to a tiered release for GPT-5.4-Cyber. On May 1 the Pentagon announced classified-network deployment agreements (IL6/IL7) with eight firms — AWS, Google, Microsoft, OpenAI, SpaceX, NVIDIA, Reflection, and Oracle — with Anthropic notably excluded over a usage-policy dispute around domestic surveillance and autonomous weapons. The list confirms a clear “sovereign AI” market that frontier-model and platform vendors are racing to occupy, and creates a meaningful commercial penalty for safety-policy stringency.

Funding rounds added to tracker

None of material size in the AI+security pure-play category surfaced in the April 28–May 4 window. Crunchbase’s top-10 weekly U.S. roundup (April 25–May 1) was led by True Anomaly (\$600M, space security/defense), with the rest concentrated in fintech, AI customer service, and developer tools; no AI cybersecurity company appeared. This is consistent with the post-RSAC pattern in 2024 and 2025, where deals announced around the conference fill March, then taper for two to three weeks. We expect activity to resume in mid-May.

Exits & M&A

2026-04-30 — Palo Alto Networks announced intent to acquire Portkey (AI Gateway). Terms undisclosed. Defines the “AI Gateway” control-plane category and integrates Portkey into Prisma AIRS.

2026-04-30 — Palo Alto Networks × Armadin partnership announced. Not an acquisition, but the first major embed of an autonomous-AI offensive-security agent inside a top-three platform’s incident-response service.

2026-04-20 (catch-up) — ServiceNow closed Armis acquisition for \$7.75B. Adds OT/IoT/medical/cloud exposure-management to the ServiceNow security platform; closes the gap between asset visibility and cyber risk.

New entrants and notable launches

Sevii launched its Cyber Swarm Defense (CSD) mode on April 28 inside its autonomous Defense and Remediation (ADR) platform. The notable design choice is commercial, not technical: CSD charges by asset protected, not by AI-token consumption, and uses Sevii’s

“Myrmidon Defense Technology” to dynamically scale agent fleets while guaranteeing sub-15-minute remediation SLOs. This addresses the single biggest objection enterprise buyers have to agentic AI security — unbounded token cost — and we expect copy-cat pricing models from Dropzone, Prophet, and Simbian over the next quarter.

On April 30, Anthropic launched Claude Security in public beta, accessible from claude.ai/security. The product is a vulnerability scanner built on Opus 4.7 that reasons over entire codebases, traces data flow, attaches confidence ratings, and ships with a multi-stage validation pipeline to suppress false positives. CrowdStrike, Palo Alto Networks, SentinelOne, Trend Micro (TrendAI), and Wiz are integrating Opus 4.7 into their platforms in parallel. This makes Anthropic both a foundation-model supplier and a direct cyber product vendor, and accelerates the “every model lab is also a security vendor” thesis we have been tracking since late 2025.

On May 1, Cisco open-sourced its Model Provenance Kit (MPK) on GitHub, with an initial fingerprint database of approximately 150 base models from 45 families and 20 publishers, accessible on Hugging Face. MPK uses architecture metadata, tokenizer structure, and weight-level signals to determine if two transformer models share a common origin. The release validates AI supply-chain integrity as a real customer problem and creates the de facto schema future commercial products will need to interoperate with.

We are adding to the Companies tab: Portkey (acquired by PANW, announced), Armadin (Mandia’s offensive-security startup, partnered with PANW), Sevii (autonomous remediation), Reflection (frontier AI lab cleared for IL6/IL7), and Armis (now part of ServiceNow).

Shutdowns / distress

No notable AI+security pure-play shutdowns or layoffs were reported in the window. The closest signal is Anthropic’s exclusion from the Pentagon’s IL6/IL7 vendor list (May 1), which is a commercial penalty rather than a corporate failure but is meaningful for any vendor with public-sector ambitions: safety-policy stringency now has a measurable revenue cost. Tech-industry-wide layoffs continue to skew toward chip manufacturing and consulting; cybersecurity remains a net hiring sector.

Themes — what shifted this week

AI Gateway is now a named category. PANW’s Portkey acquisition follows the same pattern as “Agentic Endpoint Security” with Koi: name a category, integrate the asset into Prisma AIRS, and force every other platform vendor to respond. Watch for Microsoft (Azure AI Foundry / Copilot), Google (Vertex), CrowdStrike (Charlotte), and Cisco to ship or buy a comparable gateway in the next two quarters.

Token-cost predictability becomes a feature. Sevii’s per-asset (not per-token) pricing model directly attacks the single largest objection enterprise CISOs raise about AI-SOC vendors. Expect competing vendors to follow within one to two quarters, and expect this to become a procurement-criteria checkbox.

Frontier AI labs are now cyber product vendors. Anthropic Claude Security (April 30) is a direct, branded vulnerability-scanning product, not just a model API. OpenAI’s tiered GPT-5.4-Cyber release follows the same direction. The implication for security startups: any product whose moat is “we wrote good prompts on top of a frontier model” is now competing with the frontier lab itself. Defensible wedges shift to proprietary telemetry, vertical integration, customer-tenant data residency, and policy/governance.

Sovereign AI is bifurcating along policy lines. Pentagon’s IL6/IL7 list (AWS, Google, Microsoft, OpenAI, SpaceX, NVIDIA, Reflection, Oracle) confirms a 7–8-vendor public-sector AI oligopoly. Anthropic’s exclusion creates a clear two-track market: vendors willing to ship with

broad-use clauses for defense customers, and vendors with stricter usage policies that will be limited to commercial. This reinforces our Idea #17 thesis (sovereign agentic SOC) and increases the value of model-agnostic agent runtimes.

AI supply-chain integrity moves from theoretical to operational. Cisco's open-sourced Model Provenance Kit gives defenders a usable primitive for model lineage. EU AI Act enforcement and SBOM-for-AI mandates are now arriving with practical tooling. We are adding this as a new fundable wedge (Idea #20).

Time-to-exploit collapses to hours. Anthropic's Mythos disclosure, the GitHub CVE-2026-3854 RCE, and SecurityWeek's "industrial cybercrime" analysis (April 30) all point to the same operational reality. BCG's April 29 piece argues the gap is no longer tooling but synchronization: business, IT, and security must coordinate at AI speed. This favors platform vendors with end-to-end visibility (PANW, ServiceNow, Microsoft) and squeezes specialist point products that cannot integrate across the operating model.

New ideas added to the pursue list

We are adding three ideas to the master list this week, indexed 18, 19, and 20 in the tracker.

18. Vendor-neutral AI Gateway / agent control plane

Wedge: PANW's Portkey acquisition validates the AI Gateway as the agentic-AI control plane, but locks the value into Prisma AIRS. Build an open-core, multi-cloud, multi-vendor gateway that runs in the customer's tenant and supports Anthropic, OpenAI, Google, Bedrock, Azure, and self-hosted models with consistent identity, policy, and observability. Differentiation: identity-aware semantic routing, deep MCP integration, on-prem and sovereign deployment, and multi-LLM equality. Buyers: CrowdStrike-aligned and GCP-native enterprises, FedRAMP/IL5+ regulated agencies. Risk: PANW + Microsoft + Google bundle similar gateways into existing platforms; defend by being the neutral, open option.

19. Autonomous offensive AI for the mid-market

Wedge: Armadin and PANW Unit 42 are building autonomous attack swarms for the Fortune 500. The mid-market — regional banks, mid-market SaaS, MSP-served verticals (legal, accounting, healthcare clinics) — cannot afford that, but is now in the blast radius of Mythos-class attackers. Build a self-serve continuous-validation product, sold through MSPs and cyber-insurance carriers as an underwriting input. Differentiation: pre-built per-vertical attack playbooks, channel-only GTM, and pricing tied to insurance premium discounts. Risk: channel sales cycles are slow; the answer is to seed with two cyber insurers in year one and let underwriting demand pull the channel.

20. AI model provenance and supply-chain integrity SaaS

Wedge: Cisco's Model Provenance Kit (May 1) ships the open-source primitive; enterprises pulling hundreds of models from Hugging Face need a commercial layer with attestation, signed lineage, integration with model registries (HF, Databricks, Vertex, SageMaker), continuous integrity monitoring, and GRC mappings to EU AI Act / NIST AI RMF / ISO 42001. Buyers: regulated industries (FinServ, Pharma, Defense), AI platform teams, model marketplaces. Risk: Cisco may bundle a commercial version, and Anthropic/OpenAI may ship native attestation; differentiate by being multi-vendor, multi-source, and audit/regulatory-mapped.

Index hygiene

Companies tab now has 115 entries (up from 110). Funding Rounds: 36 (no change — no qualifying rounds in window). Exits & M&A: 21 (up from 18). Startup Ideas: 20 (up from 17). Sources: 31 (up from 21).

Weekly Update — Week of 2026-04-27

This is the first weekly update following the 2026-04-20 initial survey. New entries are appended to the tracker; this section captures the narrative for the week and any change in our thesis.

Headline events

Two truly new companies entered the public landscape this week, and Palo Alto Networks formalized a category we had been tracking informally. Spectrum Security emerged from stealth on April 22 with a \$19M seed led by TechOperators; the company attacks detection engineering itself rather than the alert pile, and reports cutting detection authoring time from 121 days to under 30 minutes. Skinos Ventures — a new fund from Cato/Imperva founder Shlomo Kramer and Yishay Yovel — participated, an early signal that detection automation is now a standalone wedge separate from AI-SOC. On April 23, McLean-based Rilian raised \$17.5M in seed and seed-extension capital led by 8VC, First In, and Tamarack Global, with the UAE Cybersecurity Council / NSOC as the anchor customer for its agentic orchestration platform "Caspian." Rilian is the clearest validation we have yet seen of sovereign / nation-scale agentic security as a distinct go-to-market motion.

On the platform side, Palo Alto Networks closed its acquisition of Israeli startup Koi on April 14 (reported around \$400M) and used the deal announcement to define a new product category, "Agentic Endpoint Security" (AES), integrated into Prisma AIRS and Cortex XDR. Combined with the formal close of the Google/Wiz \$32B deal on March 26, the cloud and endpoint platform stack is now reorganizing around AI agents as first-class principals. This reinforces our March-April working hypothesis: every major incumbent is buying its way into agentic security rather than building, and the window for specialist M&A in agent identity, agent runtime, and agent data security is roughly 12 to 18 months.

Funding rounds added to tracker

2026-04-23 — Rilian: \$17.5M seed + extension; lead 8VC / First In / Tamarack Global; agentic AI for sovereign cyber + defense; UAE NSOC anchor.

2026-04-22 — Spectrum Security: \$19M seed; lead TechOperators (with Skinos Ventures, WhiteRabbit, Alumni); detection-engineering automation for AI-era SOC.

2026-04-15 — Artemis: \$70M Series A; lead Felicis (with First Round, Brightmind, Theory; angels from Demisto, Abnormal, Splunk, CrowdStrike, PANW, MSFT, Okta); AI-native defense.

2026-04-15 — Gitar: \$9M seed; lead Venrock; agents that secure code commits and pull requests.

2026-03-31 — Linx Security: \$50M Series B; lead Insight Partners (with Cyberstarts, Index); identity governance covering human, NHI, and agent identities. Total funding now \$83M.

Exits & M&A

2026-04-14 — Palo Alto Networks completed acquisition of Koi (~\$400M). Defines new "Agentic Endpoint Security" category. Integrates with Prisma AIRS and adds a new Cortex XDR module.

2026-03-26 — Google formally closed its \$32B Wiz acquisition (one year after announcement).

New entrants and notable startups added

In addition to Spectrum and Rilian, we are adding the RSAC 2026 Innovation Sandbox cohort to the tracker, which is the strongest single-event indicator of where the industry consensus on AI security is headed. Geordie AI (London) won "Most Innovative Startup" with an agent-native security and governance platform. Other top-10 finalists were Token Security, Realm Labs,

Charm Security, Clearly AI, Crash Override, Fig Security, Glide Identity, Humanix, and ZeroPath. Each finalist received a \$5M investment as part of the finalist program. Four of the ten directly target agent governance, identity, or runtime — agentic AI security is now unambiguously the dominant theme of the year.

We also backfilled a small number of late-Q1 2026 entries that were not in the initial survey: Onyx Security (Israel; agentic AI security; \$40M out-of-stealth in March), Trent AI (\$13M seed), Ray Security (\$11M seed; predictive data security for AI), and Tenex.AI (\$250M Series B; AI-enabled MDR).

Shutdowns / distress

No material AI-security shutdowns or layoffs surfaced this week. Generic Q1 2026 tech layoff data continues to skew toward chip and consulting; the most recent material event in our coverage area remains Deepwatch (November 2025).

Themes — what shifted this week

Detection engineering as a wedge. Spectrum and the early signal from Skinos Ventures imply that authoring, testing, and maintaining detection logic is a distinct buyer than "AI-SOC analyst." The buyer is a detection engineer / threat-hunting lead, not a Tier-1 SOC manager. Expect open-source competitors within 6 months built on Sigma / OCSF.

Sovereign agentic security is now fundable. Rilian + UAE NSOC is the first US-VC-backed sovereign agentic security deal we have logged. EU AI Act, NIS2, GCC localization, and ITAR / EAR are pushing customers toward on-soil deployments. Margin profile sits between defense prime and SaaS; expect a services overlay model.

"Agentic Endpoint Security" is now a category. PANW formalizing AES is a meaningful signal — once a top-three platform names a category, specialists must either get acquired before saturation or compete on agent-framework depth (LangGraph, CrewAI, AutoGen) that platform vendors cannot match.

RSAC ISB confirms agentic AI is the dominant theme. Four of ten finalists target agent governance / identity / runtime; one targets AI governance broadly; two target agentic identity. Five of ten ISB finalists touch agents in some way. This is the strongest single-event signal in the report.

New ideas added to the pursue list

We are adding two ideas to the master list this week, indexed 16 and 17 in the tracker.

16. Detection-engineering-as-code platform

Wedge: detection authoring is still a 90-120 day cycle in most enterprises. Spectrum is building a closed-SaaS attack on this. The contrarian play is open-core — a framework-led, multi-cloud detection generator for cloud-first SOCs and MSSPs, with closed-source eval and deployment plane. Distribution via Sigma / OCSF community and MSSP channel; pricing per detection authored per environment per year. Differentiation: multi-SIEM coverage day one (Splunk, Sentinel, Elastic, Snowflake) and an agent that learns from each customer's telemetry corpus.

17. Sovereign agentic SOC for regulated and public-sector markets

Wedge: GCC, EU public sector, and Allied defense customers cannot deploy US-hosted agentic SOC vendors due to data-residency, sovereignty, and ITAR / EAR rules. Rilian + UAE NSOC validated the demand. Build an air-gappable agent runtime, model-agnostic (locally hosted Llama/Qwen/sovereign clouds), region-resident operations, with FedRAMP High / IL5 / EU TI /

GCC trust mark roadmap. Distribution via local SIs (DTS, G42, Atos, Thales). Margin profile is defense-prime-adjacent; expect a 30-40% services line for the first three years.

Index hygiene

Companies tab now has 110 entries (up from 90). Funding Rounds: 36 entries. Exits & M&A: 18 entries. Startup Ideas: 17. Sources: 20.